

# Hackathon KAUST GPU 2020

Marcin Rogowski, Suha Kayum & Vincent Etienne

KAUST & Saudi Aramco

November 2020

- 1 **hpcscan**
  - Overview
  - Compilation and validation
- 2 **Test platforms**
  - Shaheen II (KAUST)
- 3 **Test Case Grid**
- 4 **Test Case Propa**
- 5 **Conclusions and next steps**
- 6 **Acknowledgements**

- 1 **hpcscan**
  - Overview
  - Compilation and validation
- 2 Test platforms
  - Shaheen II (KAUST)
- 3 Test Case Grid
- 4 Test Case Propa
- 5 Conclusions and next steps
- 6 Acknowledgements

hpcscan is a C++ code for benchmarking HPC kernels (mainly for solving PDEs with FDM)

- Simple code structure based on individual test cases
- Easy to add new test cases
- Main class is Grid: multi-dimension (1, 2 & 3D) Cartesian grid
- Hybrid MPI/OpenMP parallelism
- All configuration parameters on command line
- Support single and double precision computation
- Compilation with standard Makefile
- No external libraries
- Follows C++ Google style code

## hpcscan embeds several test cases

### Current version 1.0

- General operations on grids
- Memory operations
- MPI communication
- FD computation
- Basic wave propagator

### Possible additions for future versions

- Operations on matrices full and sparse
- FFT
- IO
- Compression

# Compilation and validation

Compiling hpcscan

go to `./build` and `make` (by default compilation with single precision float)

To compile with double precision float, make `precision=double`

Validating hpcscan

go to `./script` and `sh runValidationTests.sh`

Table: `runValidationTests.sh`<sup>1</sup>

| Machine | Compiler        | Single prec.                        | Double prec.                        |
|---------|-----------------|-------------------------------------|-------------------------------------|
| Mars    | g++ 9.3.0       | 764 PASS / 0 FAIL / 0 ERR / 20 WARN | 764 PASS / 0 FAIL / 0 ERR / 20 WARN |
| Shaheen | icpc 19.0.5.281 | 764 PASS / 0 FAIL / 0 ERR / 20 WARN | 764 PASS / 0 FAIL / 0 ERR / 20 WARN |

Numbers can differ due to availability of features depending on the platforms

---

<sup>1</sup>Updated Nov 25, 2020

- 1 hpcscan
  - Overview
  - Compilation and validation
- 2 Test platforms
  - Shaheen II (KAUST)
- 3 Test Case Grid
- 4 Test Case Propa
- 5 Conclusions and next steps
- 6 Acknowledgements

# Test platform - Shaheen II (KAUST)

## Machine Shaheen II / Cray XC40

- Computing nodes Intel Haswell 2.3 Ghz dual socket (16 cores / socket)
- RAM 128 GB with Peak memory BW 136.5 GB/s
- Peak performance Single Prec. 2.36 TFLOP/s / Double Prec. 1.18 TFLOP/s
- Interconnect Cray Aries with Dragonfly topology
  - 60 GB/s optical links between groups
  - 8.5 GB/s copper links between chassis
  - 3.5 GB/s backplane within a chassis
  - 5 GB/s PCIe from node to Aries router





- 1 hpcscan
  - Overview
  - Compilation and validation
- 2 Test platforms
  - Shaheen II (KAUST)
- 3 Test Case Grid
- 4 Test Case Propa
- 5 Conclusions and next steps
- 6 Acknowledgements

# Test Case Grid - Description

- Fill grid ( $W = \text{coef}$ )
- L1 error between grid  $W$  and  $R$
- Get min. grid  $W$
- Get max. grid  $W$
- Update pressure  $W = 2*U - W + C*L$  (used in propagator)
- Medium Grid size 4 GB ( $1000 \times 1000 \times 1000$  points)

# Test Case Grid - Results

- CPU (Intel, no way to get AMD) and GPU kernels (V100)

Table: Bandwidth GB/s <sup>2</sup>

| Machine     | Algo.    | Fill | L1 err. | Get max. | Get min. | Update Pres. |
|-------------|----------|------|---------|----------|----------|--------------|
| Shaheen CPU | Baseline | 54   | 124     | 126      | 126      | 120          |
| Ibex CPU    | Baseline | 92   | 217     | 224      | 224      | 198          |
| Ibex GPU1   | Cuda     | 497  | 274     | 700/221  | 700/221  | 673          |
| Ibex GPU2   | OpenAcc  | 270  | -       | -        | -        | -            |
| Ibex GPU3   | OpenMP   | -    | -       | -        | -        | -            |

Table: Bandwidth GPoints/s

| Machine     | Algo.    | Fill | L1 err. | Get max. | Get min. | Update Pres. |
|-------------|----------|------|---------|----------|----------|--------------|
| Shaheen CPU | Baseline | 13.5 | 15.5    | 31.5     | 31.5     | 6.0          |
| Ibex CPU    | Baseline | 23.0 | 27.1    | 56.0     | 56.0     | 9.9          |
| Ibex GPU1   | Cuda     | 124  | 34.5    | 175/55   | 175/55   | 33.7         |
| Ibex GPU2   | OpenAcc  | 67   | -       | -        | -        | -            |
| Ibex GPU3   | OpenMP   | -    | -       | -        | -        | -            |

Reproduce results with `./hackathonTestCases/testCase_Grid/`

<sup>2</sup> Updated Dec 1, 2020

## Machine: Shaheen

- L1 Err., Get Min & Max: 125 GB/s close to peak BW (92 % Peak Mem. BW)
- Low perf for Fill: 54-58 GB/s (40-43 % Peak Mem. BW)
- Max Err. 72-91 GB/s (53-67 % Peak Mem. BW)
- Pressure update 6 GPoint/s (120 GB/s, 88 % Peak Mem. BW)


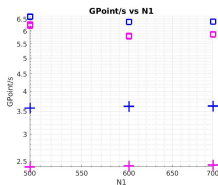

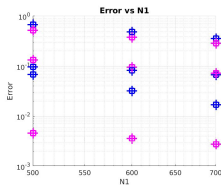
- 1 hpcscan
  - Overview
  - Compilation and validation
- 2 Test platforms
  - Shaheen II (KAUST)
- 3 Test Case Grid
- 4 Test Case Propa**
- 5 Conclusions and next steps
- 6 Acknowledgements

# Test Case Propa - Description

- Seismic wave propagator
- 2nd order acoustic wave equation
- Time-domain Finite-difference
- Various FD order in space
- 2nd FD order in time
- Various grid size and time steps
- Total 18 configurations
- Comparison against analytical solution (Eigen mode)

# Test Case Propa - Results

- **Machine: shaheen**
- 1 node / 32 threads
- CPU Baseline & cache blocking kernels <sup>3</sup>
- FD: Black O2, Blue O4, Pink O8, Red O12 / Square=CPU CacheBlk, Cross=Baseline



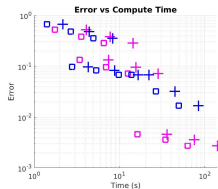



Reproduce results with `./hackathonTestCases/testCase_Propa/runShaheen.sh`

<sup>3</sup> Updated Nov 28, 2020

# Test Case Propa - Results

- **Machine: Ibex Cpu**
- 1 node / 32 threads
- CPU Baseline & cache blocking kernels <sup>4</sup>
- FD: Black O2, Blue O4, Pink O8, Red O12 / Square=CPU CacheBlk, Cross=Baseline



Reproduce results with `./hackathonTestCases/testCase_Propa/runIbexCpu.sh`

<sup>4</sup> Updated Dec 1, 2020

Marcin, Suha & Vincent

Hackathon KAUST GPU 2020

Nov 2020

16 / 21



Table: Best total time for 18 configurations <sup>5</sup>

| Machine     | Algo.    | Time (s) | Speedup |
|-------------|----------|----------|---------|
| Shaheen CPU | Baseline | 826      | 0.5     |
| Shaheen CPU | CacheBlk | 388      | 1.0     |
| Ibex CPU    | Baseline | 524      | 0.7     |
| Ibex CPU    | CacheBlk | 265      | 1.5     |
| Ibex GPU1   | Cuda     | XXX      | Y.Y     |
| Ibex GPU2   | OpenACC  | XXX      | Y.Y     |
| Ibex GPU3   | OpenMP   | XXX      | Y.Y     |

---

<sup>5</sup> Updated Dec 1, 2020

- 1 hpcscan
  - Overview
  - Compilation and validation
- 2 Test platforms
  - Shaheen II (KAUST)
- 3 Test Case Grid
- 4 Test Case Propa
- 5 Conclusions and next steps
- 6 Acknowledgements

# Conclusions and next steps

TO DO

- 1 hpcscan
  - Overview
  - Compilation and validation
- 2 Test platforms
  - Shaheen II (KAUST)
- 3 Test Case Grid
- 4 Test Case Propa
- 5 Conclusions and next steps
- 6 Acknowledgements

# Acknowledgements

- KAUST ECRC and KSL for access and support on Shaheen II & Ibex