



# Performance Characterization of a Vector Architecture for Seismic Applications

V. Etienne, A. Momin, EXPEC ARC, Saudi Aramco, Dhahran, Saudi Arabia

L. Gatineau, NEC Deutschland GmbH

S. Momose, NEC Japan

5th EAGE workshop on High Performance Computing for Upstream

6-8 September 2021



## Outline

### Introduction

Overview of this work and expected outcomes

### Seismic Exploration and High Performance Computing

Combining the best of both worlds

### Performance Characterization of a Vector Engine

Comparative benchmarks with standard architectures

### Conclusions

# Introduction

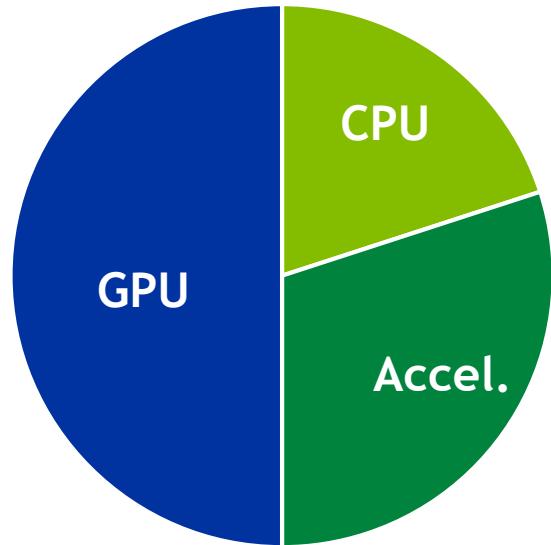
1

# Introduction

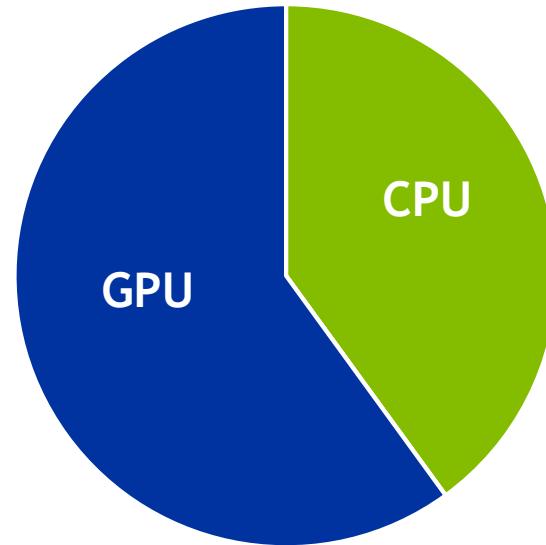
## HPC for the Oil and Gas Industry

source [top500.org](http://top500.org)

Top 10 - June 2018



Top 10 - June 2021



3 years ago

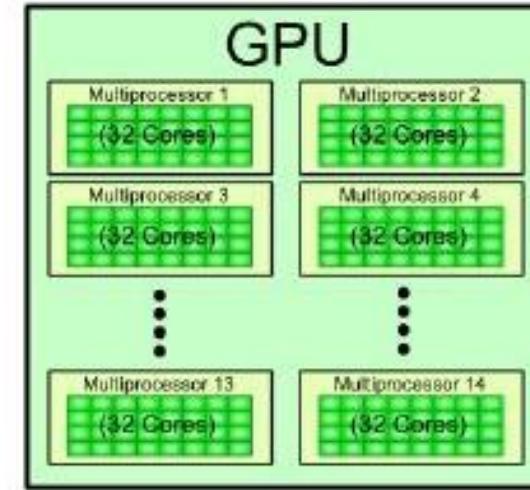
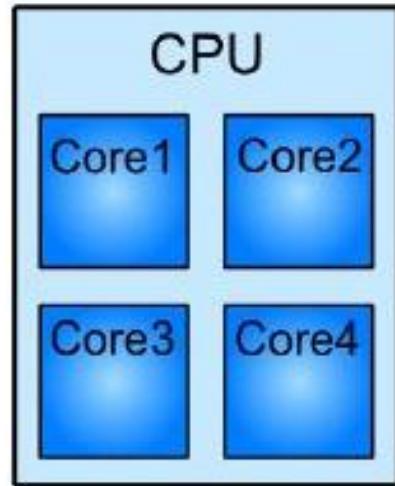
- ➊ #1 machine GPU based (0.19 EFlop/s)
- ➋ 50% machines are GPU based
- ➌ #13 ENI (GPU), #30 TOTAL (CPU),  
#35 PGS (CPU), 49 ENI (GPU)

Today

- ➊ #1 machine CPU based (0.54 EFlop/s)
- ➋ 60% machines are GPU based
- ➌ #9 ENI (GPU), #11 Saudi Aramco (GPU),  
#19 Saudi Aramco (CPU), #21 Total (GPU)

# Introduction

## Standard HPC Architectures



### Central Processing Unit

- Few tens of cores with extended capacity and performance
- Large main memory
- Large cache memory / core
- Low memory bandwidth
- Standard programming languages

Current trend: Increase memory bandwidth

### Graphics Processing Unit

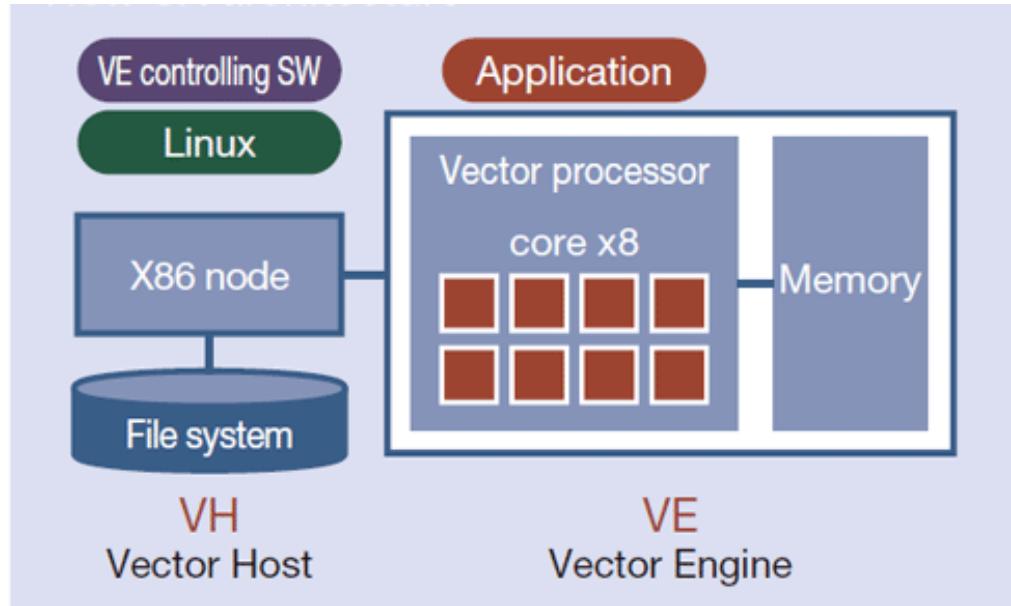
- Massive parallel device, with thousands of cores executing simple operations
- Limited main memory
- Limited cache memory / core
- High memory bandwidth
- Specific programming language

Current trend: Adopt standard languages

## Vector Architecture

# Introduction

### NEC SX-Aurora Vector Engine

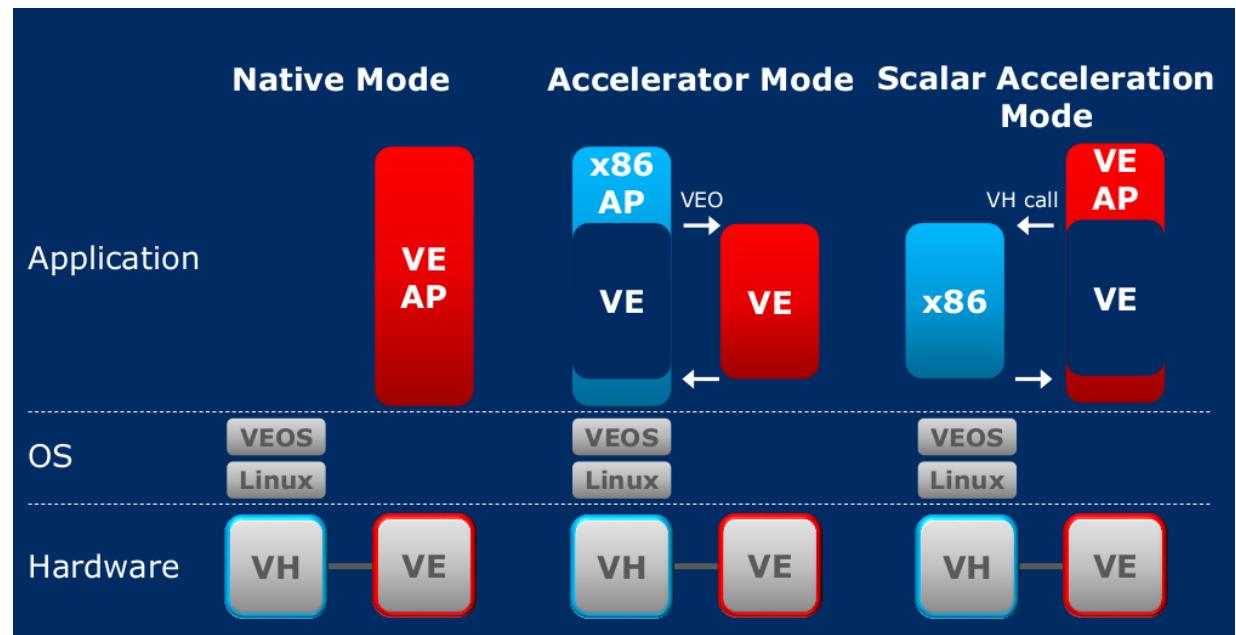


3 modes supported  
**Native (this work)**  
Accelerator  
Scalar accelerator

#### Vector Engine

- Few cores with extended capacity & performance
- Long SIMD vector length
- Limited main memory
- Large cache memory / core
- High memory bandwidth
- Standard programming languages

Combines CPU and GPU strengths



# Introduction

## Expected Outcomes of this Work

**Our main objective: a quantitative comparison between CPU, GPU and VE**

Considering the specificities of our applications, we evaluate:

 **Computation efficiency > Performance & run time**

 **Energy efficiency > Electric consumption**

 **Usability > Effort to implement & optimize algorithms**

In an objective way



With robust protocols and metrics



With an open source project available to the community



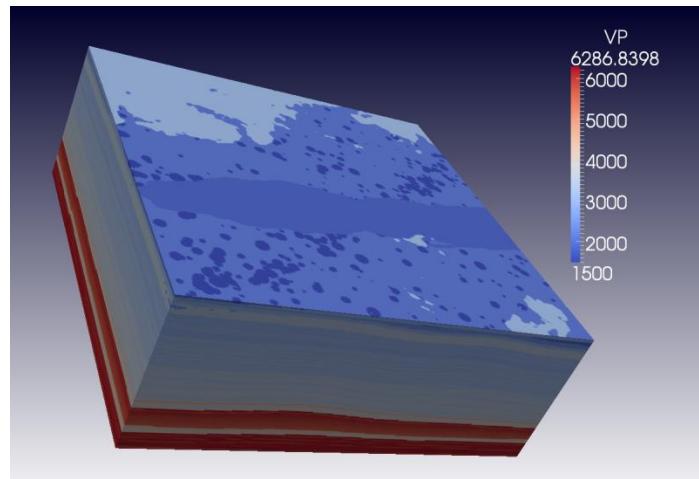
# Seismic Exploration & HPC

2

# Seismic Applications

## Seismic and HPC

In O&G  
2 main applications  
Reservoir simulation  
Seismic imaging &  
inversion

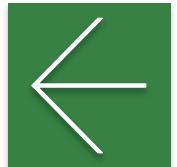


Model

Seismic modeling  
Forward problem

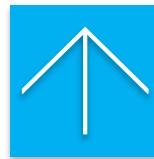


Seismic inversion  
Inverse problem

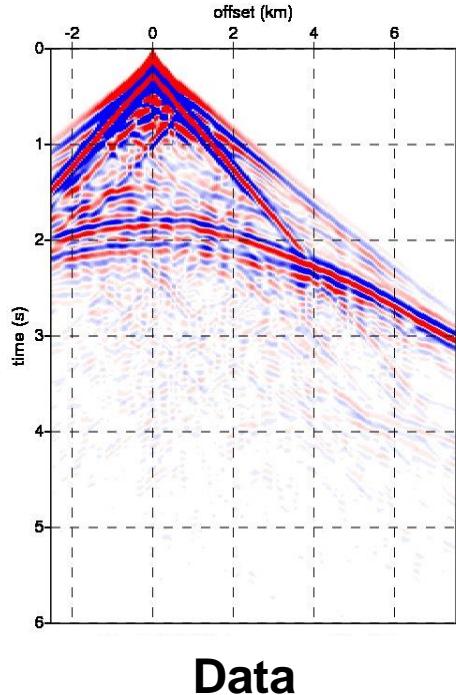


Compute intensive  
applications that require  
HPC systems

Ex. Elastic modeling in  
large scale survey  
about 100 EFlop



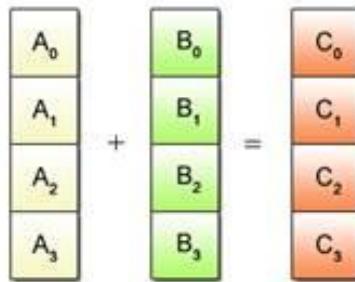
HPC



## Seismic and HPC

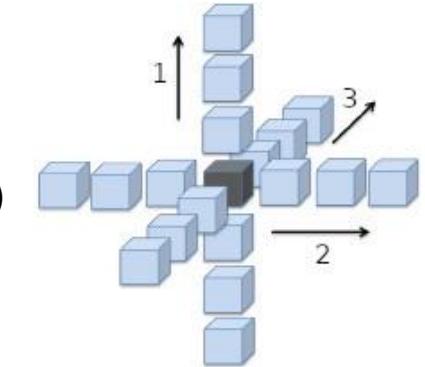
### Memory Bandwidth

- Seismic applications performs lots of operations on structured **grids**
- Most of our computational kernels are **memory bounded**



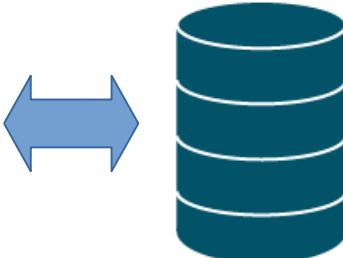
### Cache size & efficiency

- Workhorse for seismic: **Finite-difference**
- Large amount of **(in-cache) data reuse**
- Proficiency of compiler vs programmer



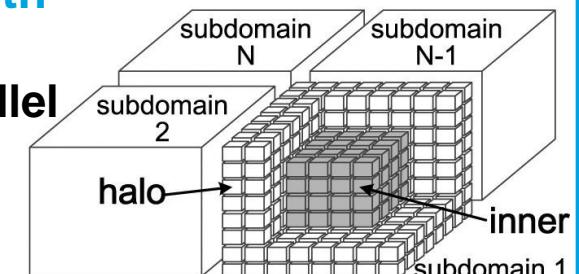
### Storage bandwidth

- **Large amount of data** written to and read from disk
- Possibly associated with **compression** algorithm



### Interconnect bandwidth

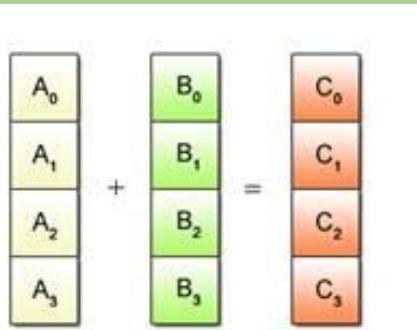
- Our applications are **embarrassingly parallel**
- Interconnect mostly used for **subdomain decomposition** (halo exchange)



## Seismic and HPC

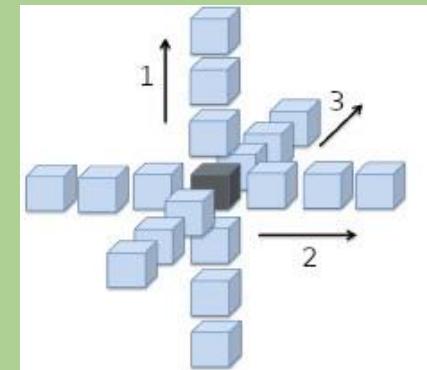
### Memory Bandwidth

- Seismic applications performs lots of operations on structured **grids**
- Most of our computational kernels are **memory bounded**



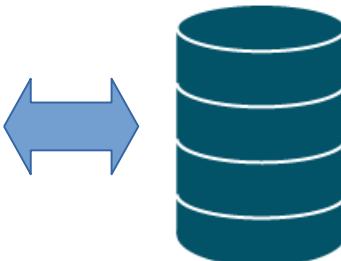
### Cache size & efficiency

- Workhorse for seismic: **Finite-difference**
- Large amount of **(in-cache) data reuse**
- Proficiency of compiler vs programmer



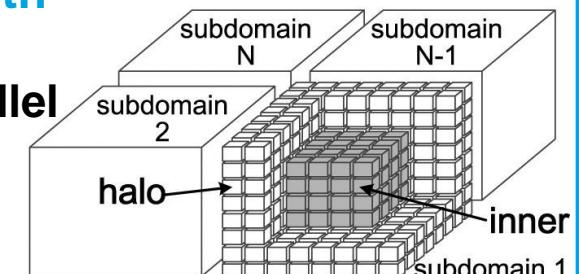
### Storage bandwidth

- **Large amount of data** written to and read from disk
- Possibly associated with **compression** algorithm



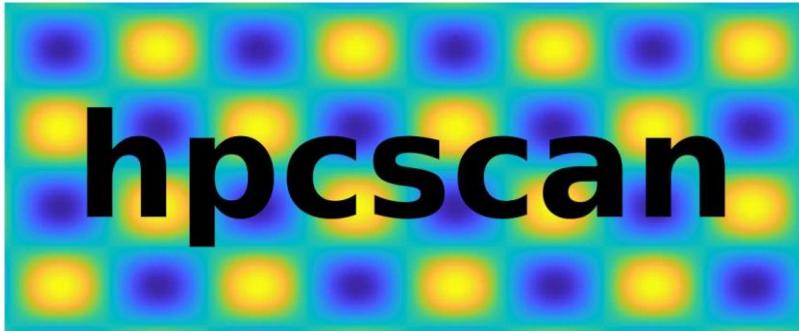
### Interconnect bandwidth

- Our applications are **embarrassingly parallel**
- Interconnect mostly used for **subdomain decomposition** (halo exchange)



In this work we will focus on the architecture-related KPIs (memory/cache & flops)

## Seismic and HPC



### HPCscan benchmark

C++ code for benchmarking HPC kernels  
with a focus on Finite Difference

Multi-architectures (CPU, GPU & VE)  
Built-in performance, accuracy and energy measurements

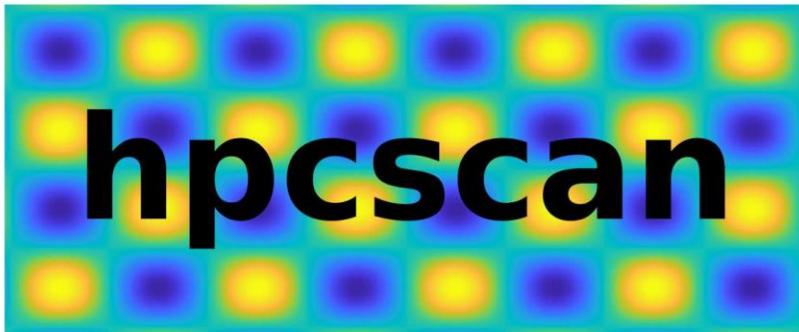
Available on GitHub  
No confidentiality: ideal for external collaborations

**Main concept: not focus on full applications (too complex)**  
**Instead, perform a 'scanning' of most sensitive kernels**

Simple to use and customize

```
mpirun -n 1 hpcscan -testMode Baseline -testCase Memory -dim 3 -n1 1000 -n2 1000 -n3 1000
```

# Seismic and HPC



## HPCscan benchmark

C++ code for benchmarking HPC kernels  
with a focus on Finite Difference

Multi-architectures (CPU, GPU & VE)  
Built-in performance, accuracy and energy measurements

Available on GitHub

No confidentiality: ideal for external collaborations

**Main concept: not focus on full applications (too complex)**  
**Instead, perform a 'scanning' of most sensitive kernels**

Simple to use and customize

```
mpirun -n 1 hpcscan -testMode Baseline -testCase Memory -dim 3 -n1 1000 -n2 1000 -n3 1000
```

- Baseline
- CacheBlk
- CUDA, CUDA\_Opt
- DPC++
- HIP, HIP\_Opt
- NEC, NEC\_SCA

- Memory
- Grid
- FD\_D2
- Comm
- Propa

# Performance characterization

3

## Test case: Memory

# Performance

Hardware specifications

**Representative products  
available on the market**

Spec.	CPU	GPU	VE
# cores	2 x 38	5 120	8
Memory type	DDR4	HBM2	HBM2
Capacity GB	256	32	48
Peak GB/s	410	900	1351

Measured memory bandwidth

Case	CPU		GPU		VE	
	GB/s	% peak	GB/s	% peak	GB/s	% peak
Fill	200	49	725	81	1060	78
Copy	300	73	616	68	1085	80
Add	301	73	694	77	1083	80
Mul	300	73	695	77	1084	80
AddUpd	294	72	691	77	1083	80

Architectures reach 70-80 % of their peak bandwidth  
 Measured bandwidth is a key indicator to assess performance of kernels

## Test case FD\_D2: FD Operators

### Performance

Starting point & Reference:  
CPU Baseline

Axis 1 & 2: High data reuse

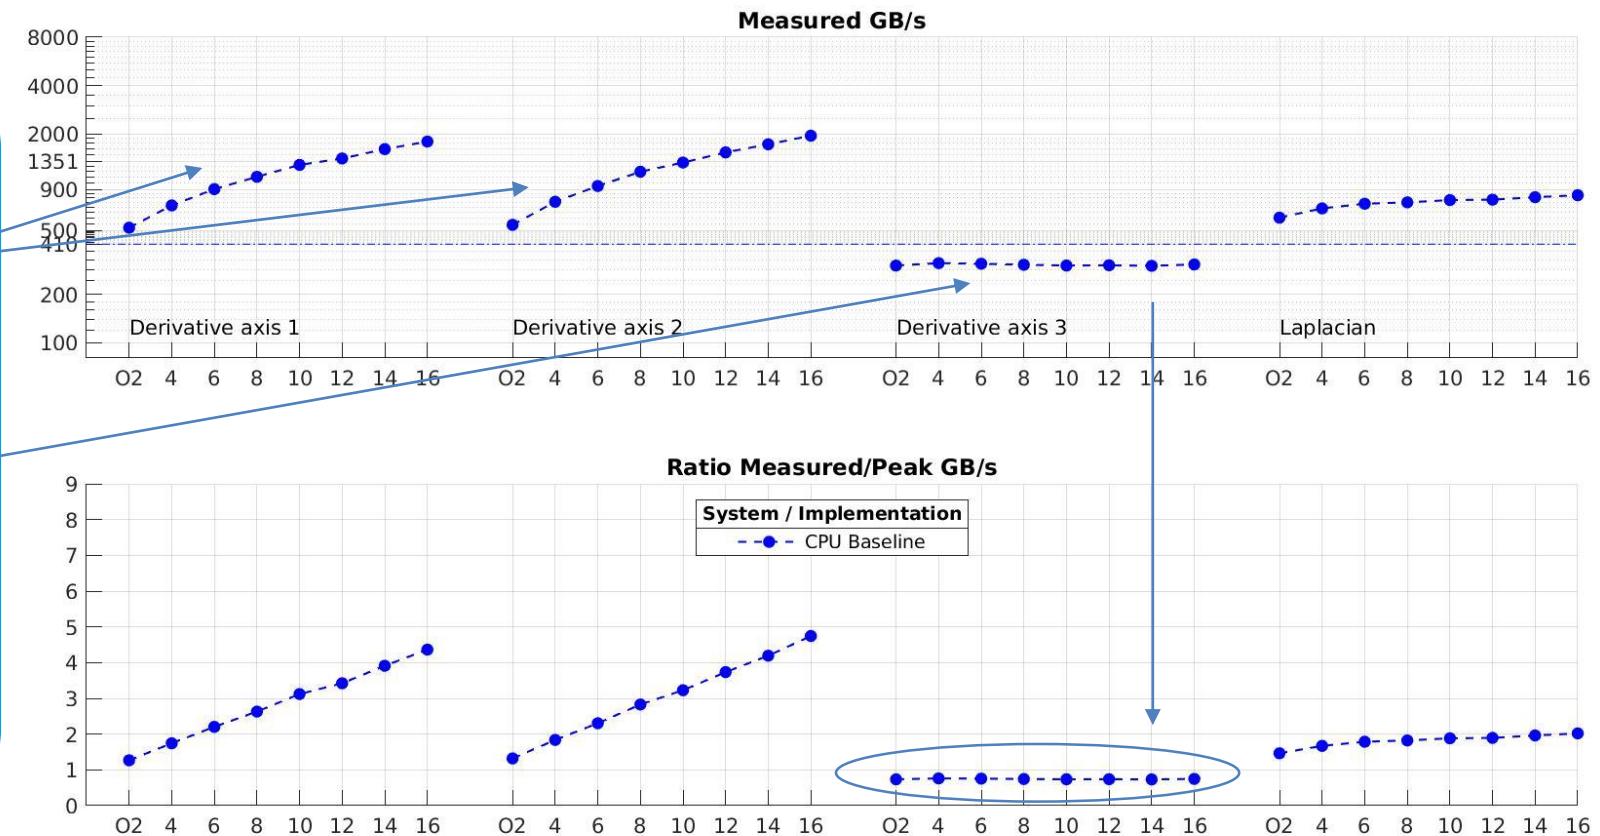
Axis 3: Low data reuse  
Large stride between points

►Characteristic pathology

No cache effect

Limited by Memory BW

Laplacian impacted by axis 3



Efficiency increases with FD order except Axis 3 and Laplacian  
How to treat pathology on Axis 3?

## Test case FD\_D2: FD Operators

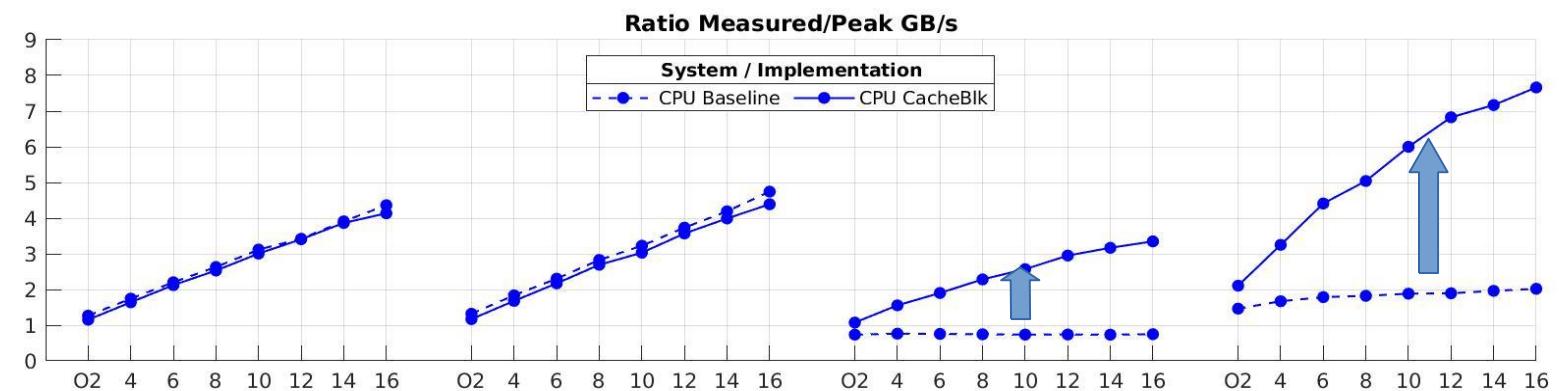
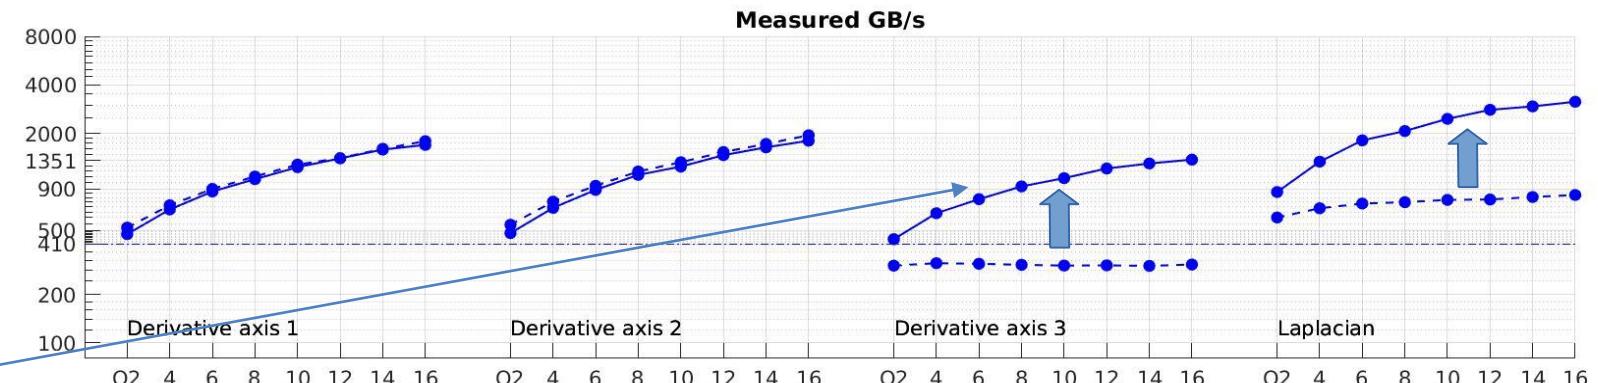
# Performance

### CPU Cache Blocking

Divide grid into small blocks  
Force data to remain in cache

Axis 3 & Laplacian  
► Considerable benefits

Implementation: easy  
Tuning: finding optimal block size is not straight-forward



Efficiency increases with FD order for all orders  
Important effect of cache blocking for the Laplacian

## Test case FD\_D2: FD Operators

# Performance

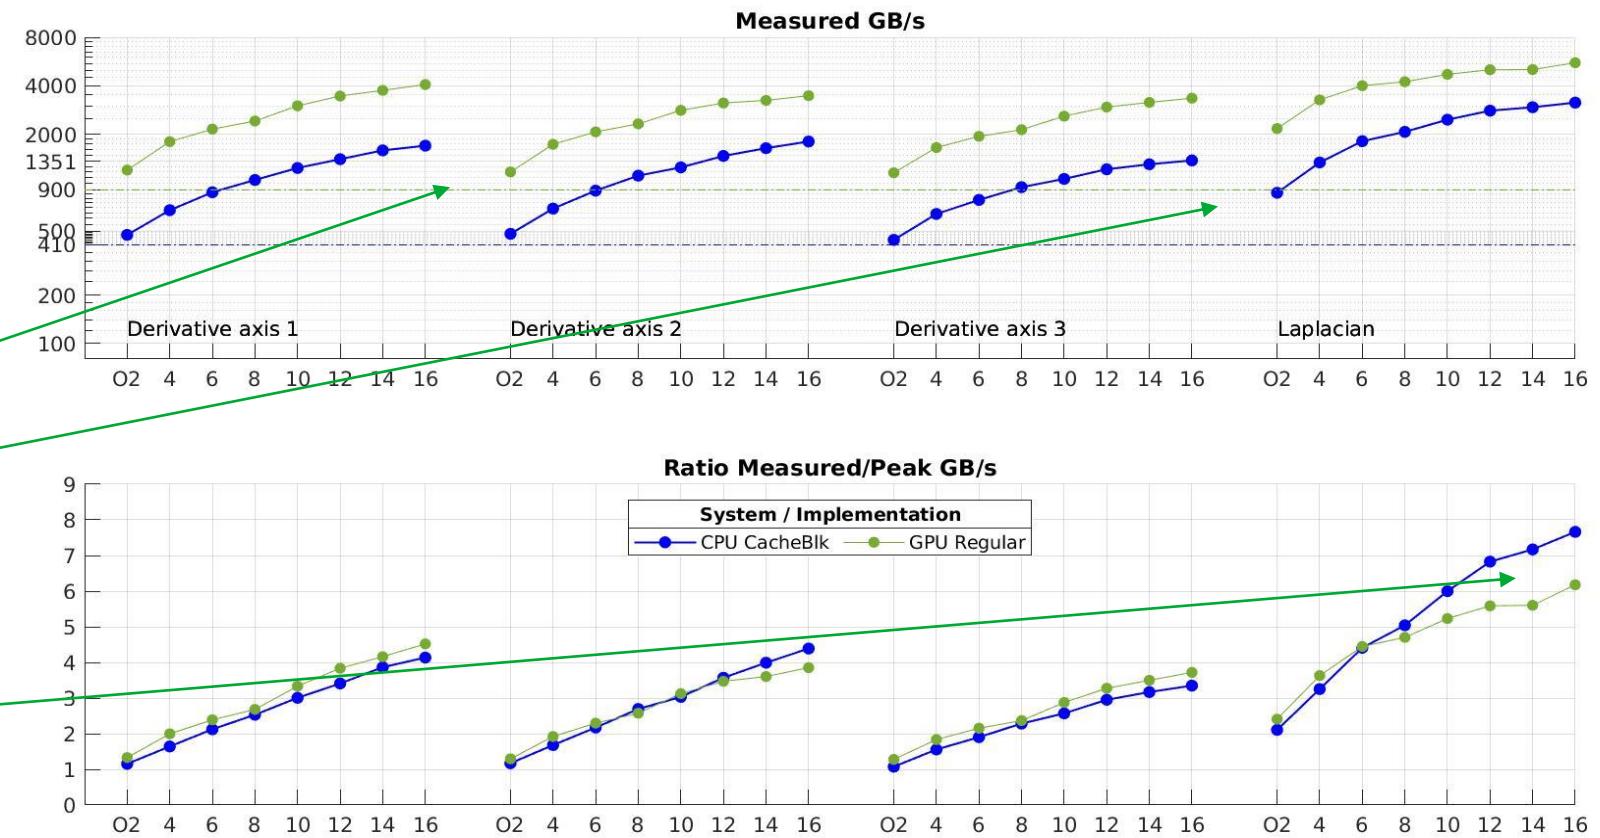
**GPU Regular implementation**  
**Constructor language**  
**Grid divided into 3D thread blocks**

Axis 1 & 2: High data reuse

Axis 3 & Laplacian: Also good

Similar relative efficiency  
between CPU & GPU

Less efficient for large FD  
orders



GPU exhibits good performances for all operators  
Can we get better?

## Test case FD\_D2: FD Operators

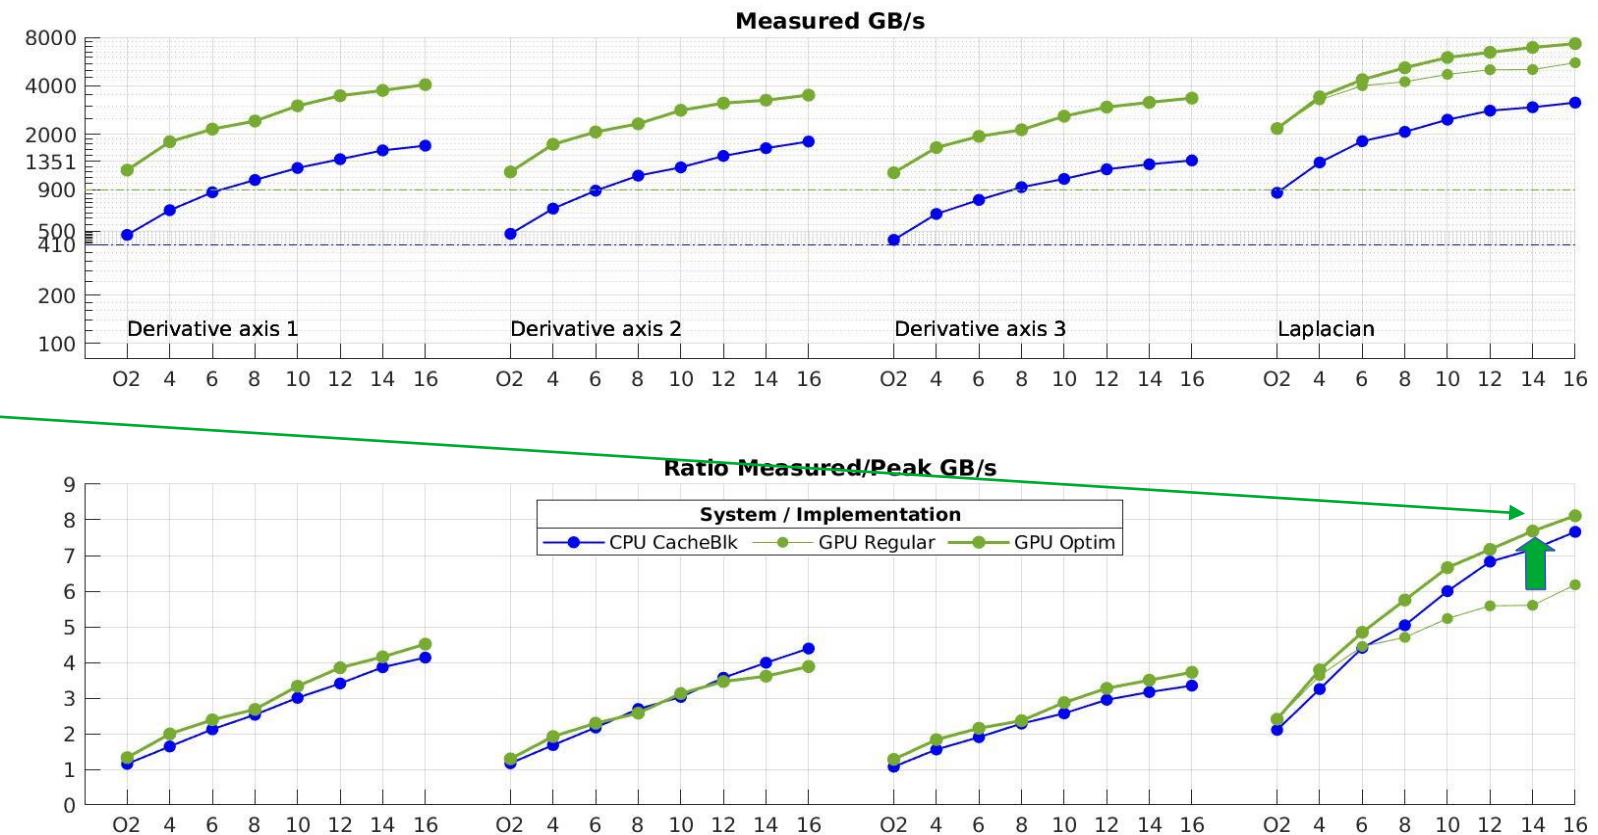
# Performance

**GPU Optim. implementation**  
**Constructor language**  
**2D ‘cache blocking’**  
**data prefetching shared memory (slice N1 x N2)**

Laplacian efficiency increased

Implementation: demanding  
Tuning: finding optimal block size is not straight-forward

Note: implemented only for Laplacian



Similar relative efficiency between CPU & GPU  
GPU naturally benefits from higher memory BW (x2.2)

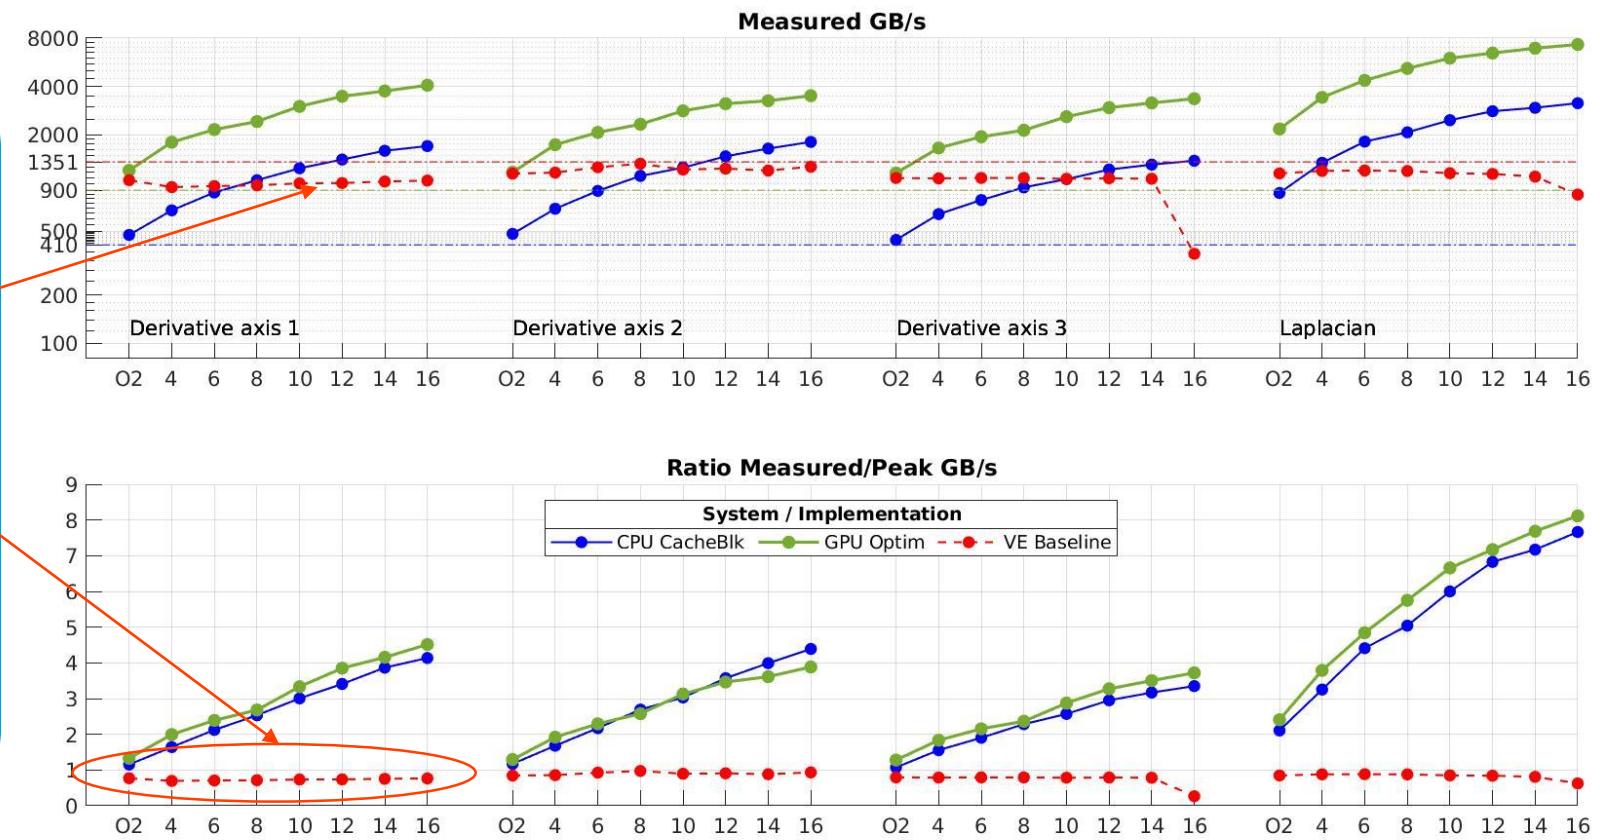
## Test case FD\_D2: FD Operators

### Performance

VE Baseline  
hpcscan simply compiled  
as is with nc++

All operators: Low data reuse

Limited by Memory BW



With a simple compilation, FD operators run at the Memory BW of VE  
There is no free lunch, let us do some effort!

## Test case FD\_D2: FD Operators

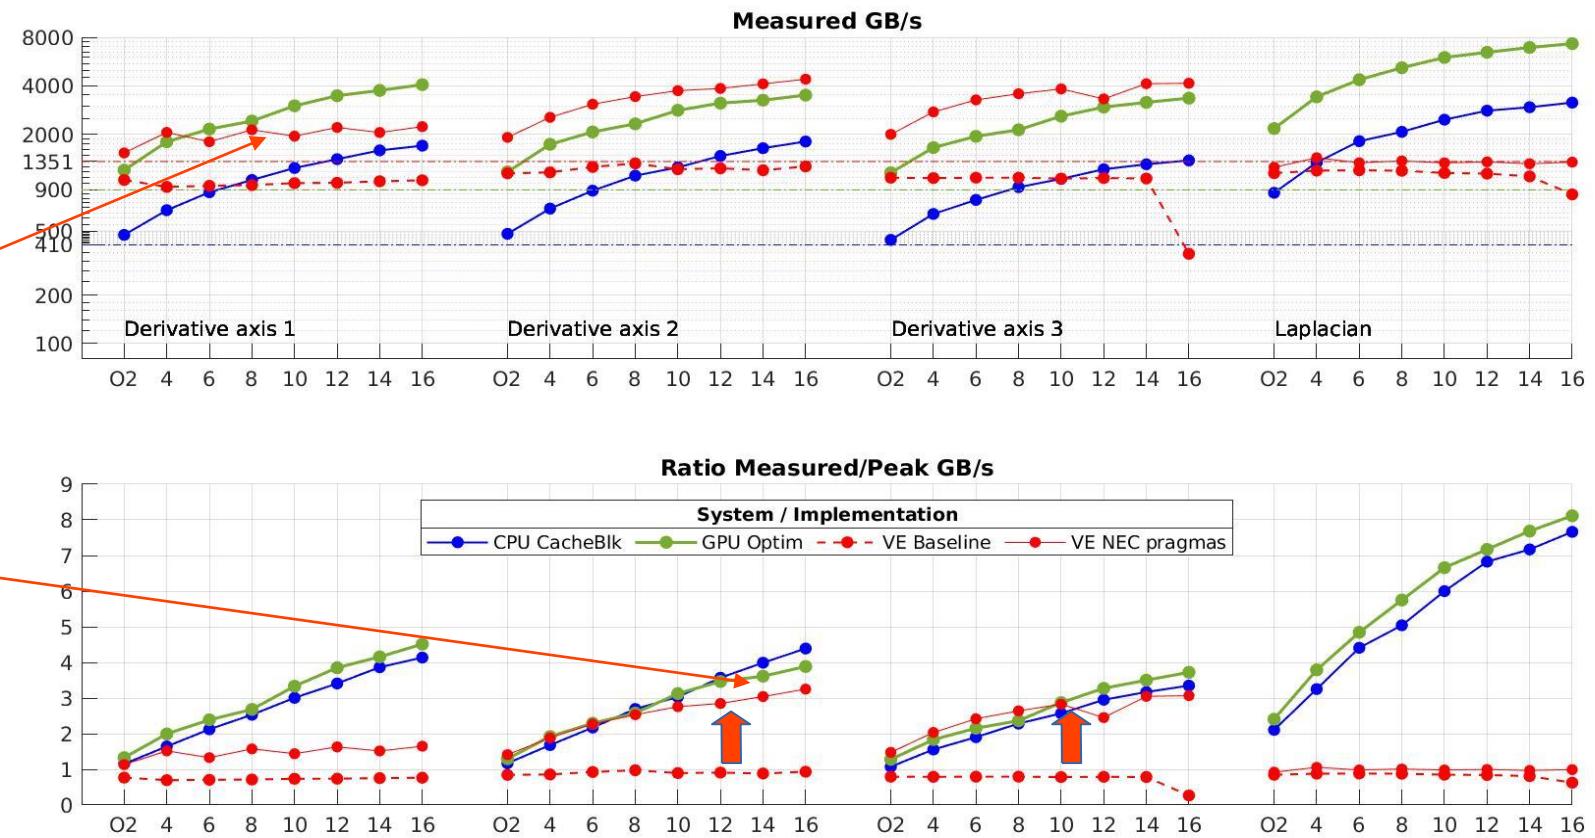
# Performance

VE with compiler directives  
NEC pragmas added to  
Baseline implementation

Axis 1: perf. Increase x2  
#pragma packed\_stencil  
vector 512 SP flops (instead  
of 256)

Axis 2 & 3: high data reuse  
#pragma outerloop\_unroll  
keep stencil data in registers

Slight improvements for the  
Laplacian



High performance for Axis 2 & 3  
What about the Laplacian?

## Test case FD\_D2: FD Operators

### Performance

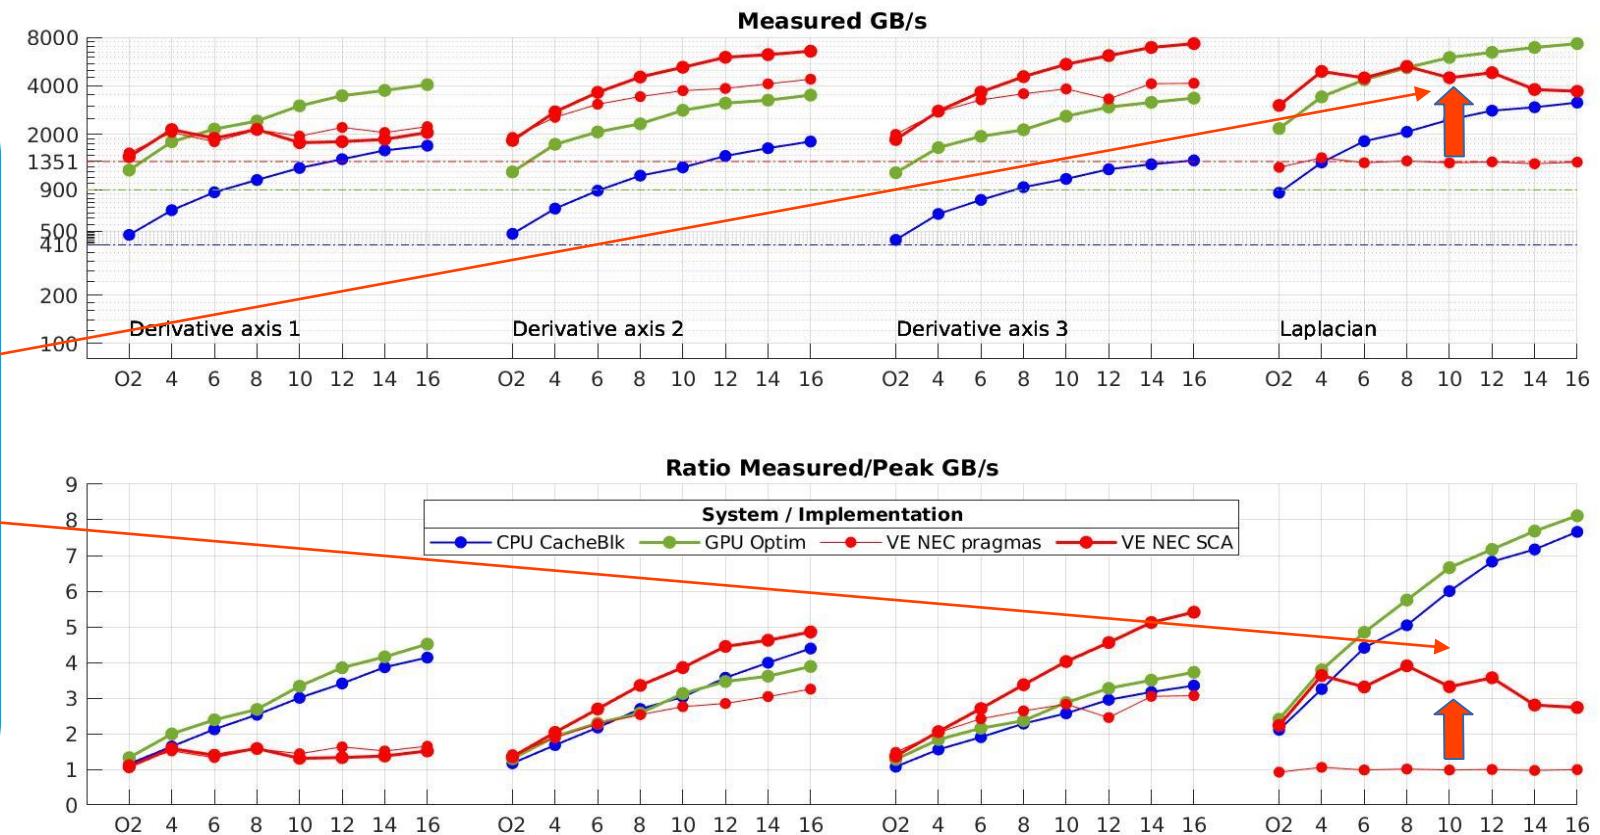
VE with NEC Math. Library  
Stencil Code Accelerator  
(SCA)  
→ Talk of S. Momose

Large improvements for Axis 2 & 3, and Laplacian

Less efficient for high orders

Bottle-neck: Axis 1

Implementation: no effort, all done by the library



**Limitation of SCA: Split computation for the acoustic propagator**  
**2 steps: 1<sup>st</sup> Laplacian computation and 2<sup>nd</sup> pressure wavefield update**

## Test case FD\_D2: FD Operators

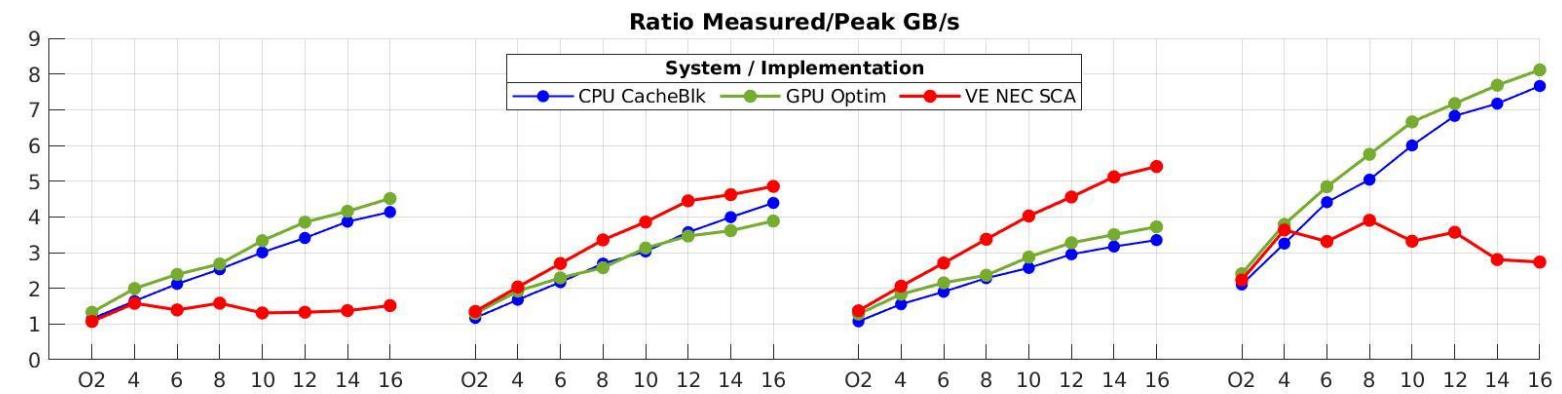
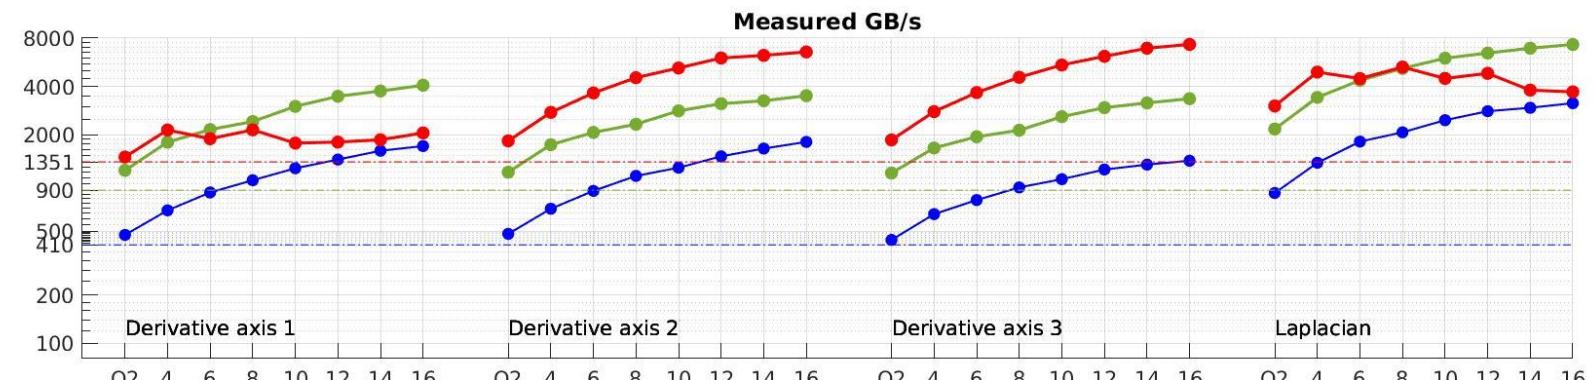
# Performance

Summary of optimal implementation for each architecture

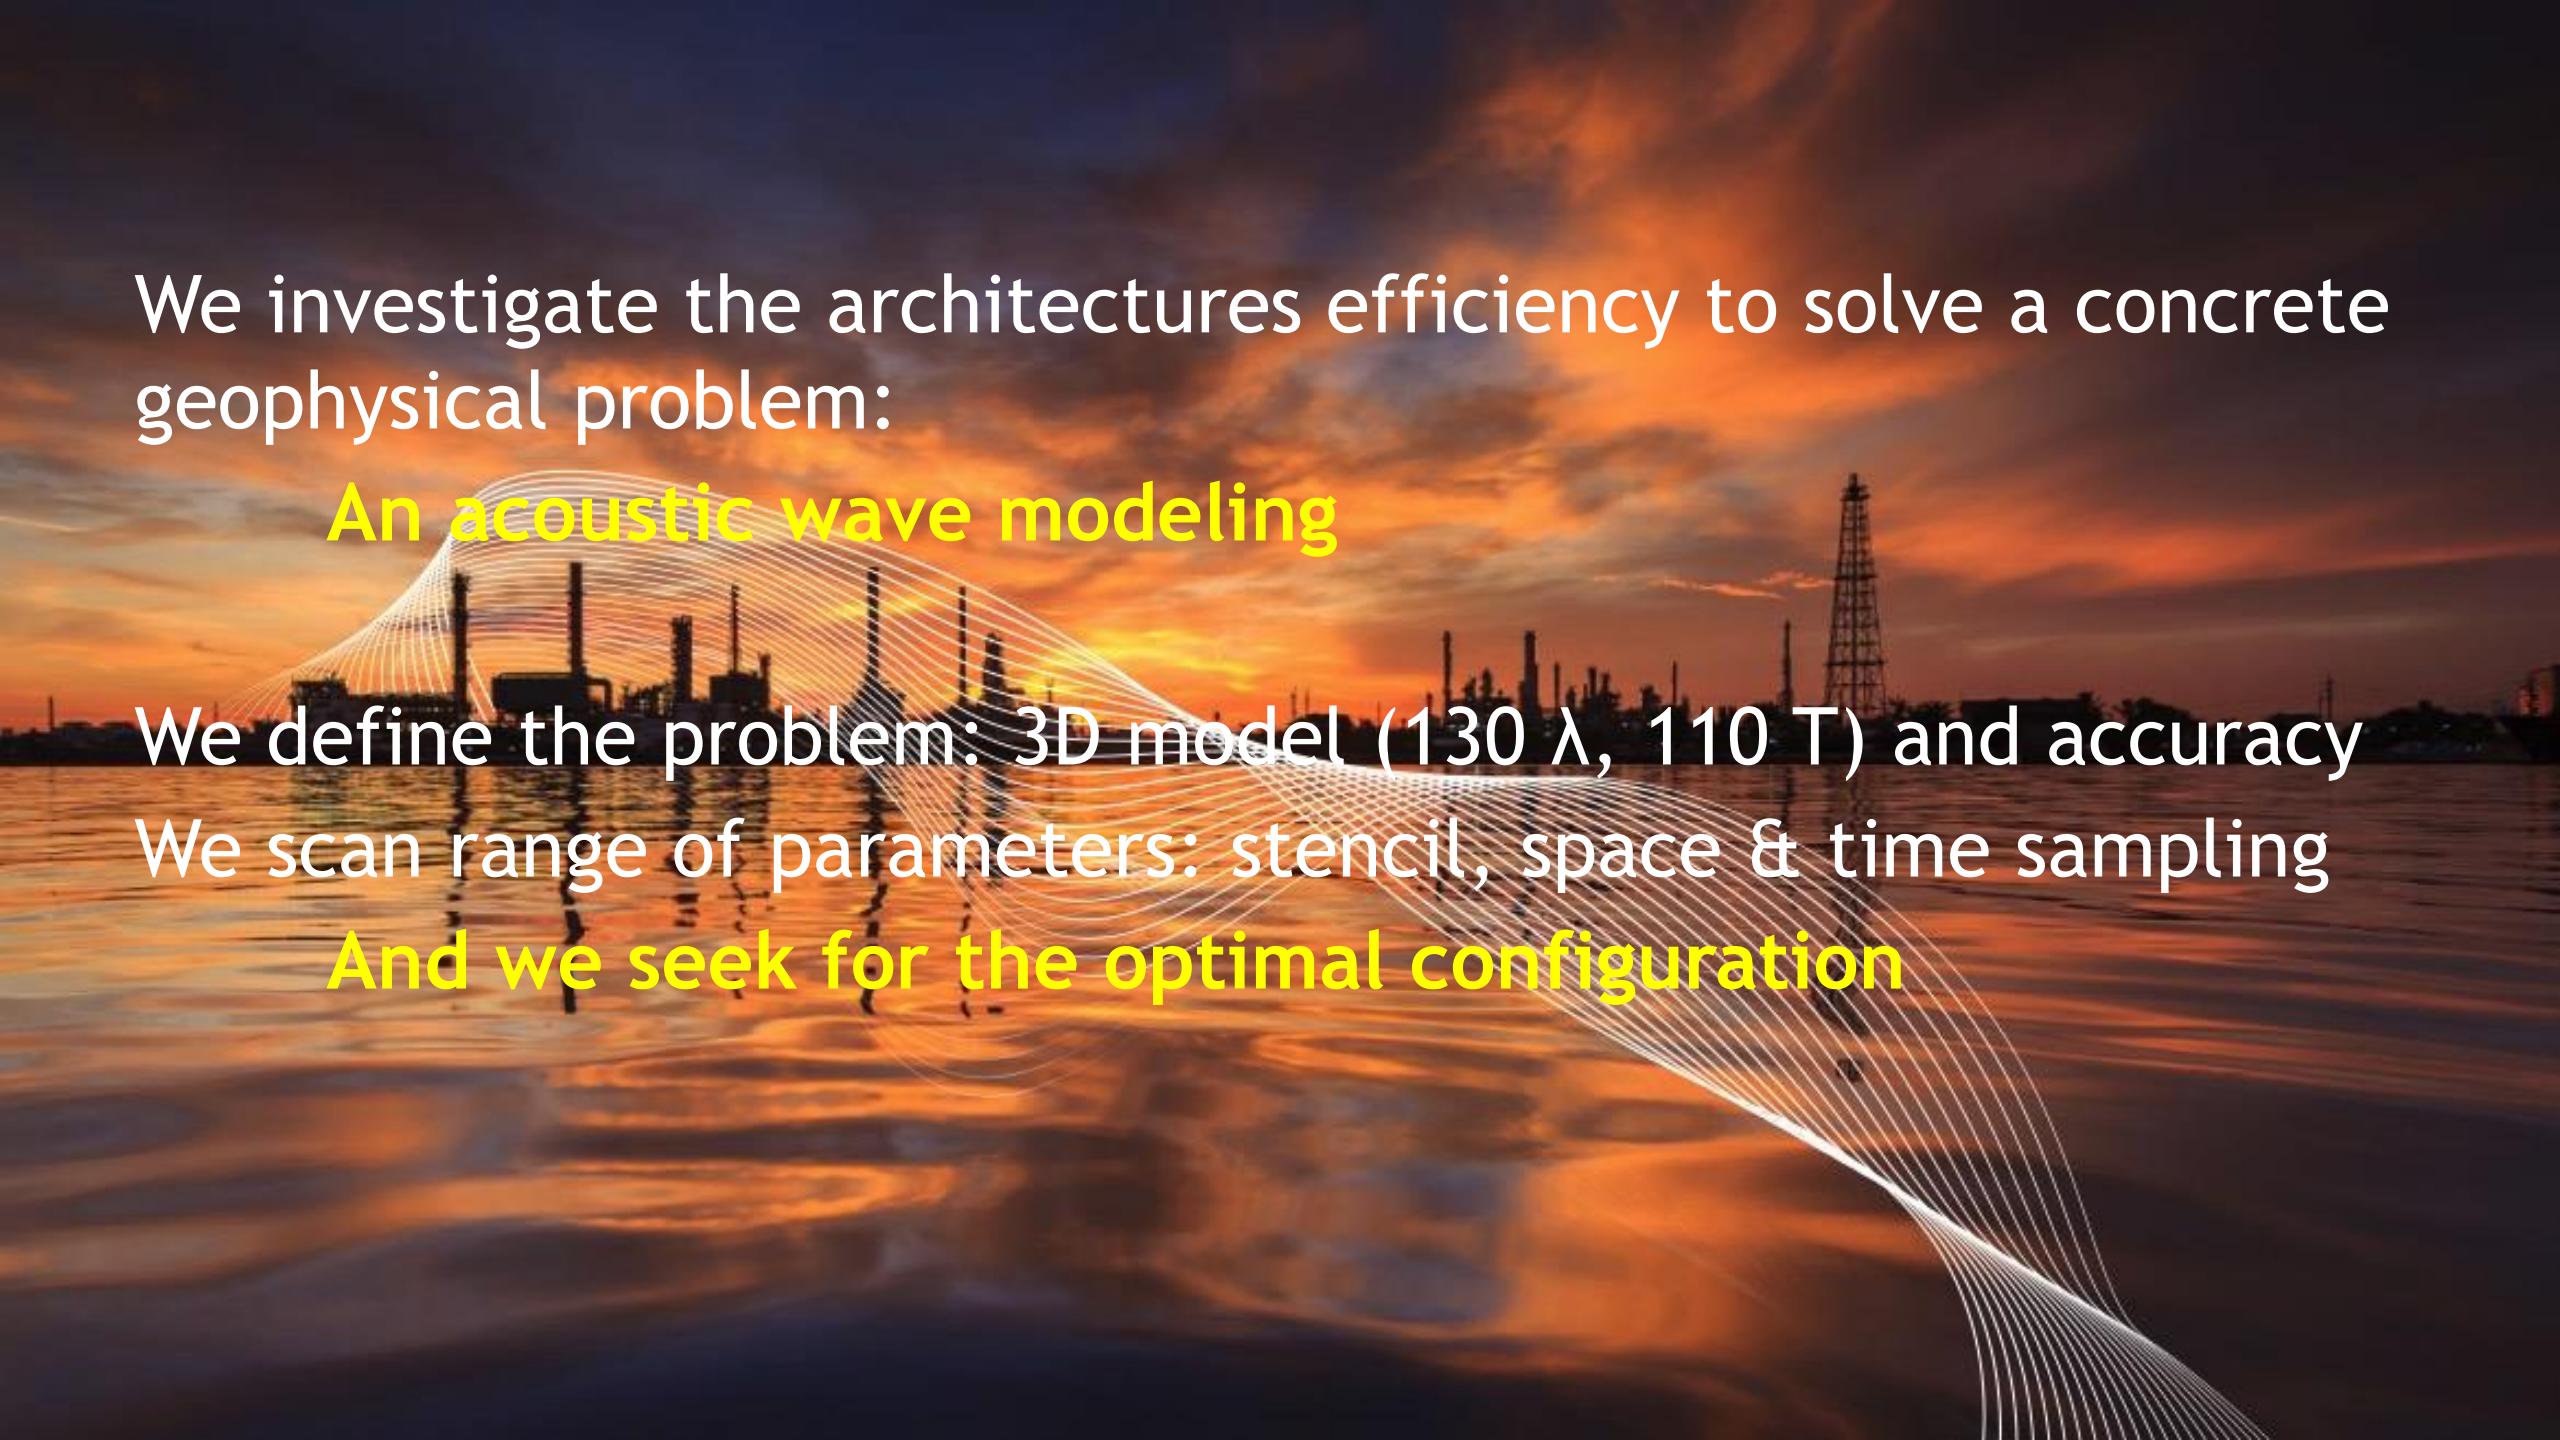
CPU: Cache blocking

GPU: Data prefetch in local (shared) memory

VE: NEC SCA library



Each architecture has its specific performance ‘signature’  
How to select the optimum?

The background of the slide features a photograph of a suspension bridge at sunset. The sky is filled with warm, orange, and yellow clouds. In the foreground, the calm water reflects these colors. Several thin, white, curved lines representing wave patterns or light rays fan out from behind the bridge's towers, creating a sense of motion and depth.

We investigate the architectures efficiency to solve a concrete geophysical problem:

### **An acoustic wave modeling**

We define the problem: 3D model ( $130 \lambda$ ,  $110 T$ ) and accuracy

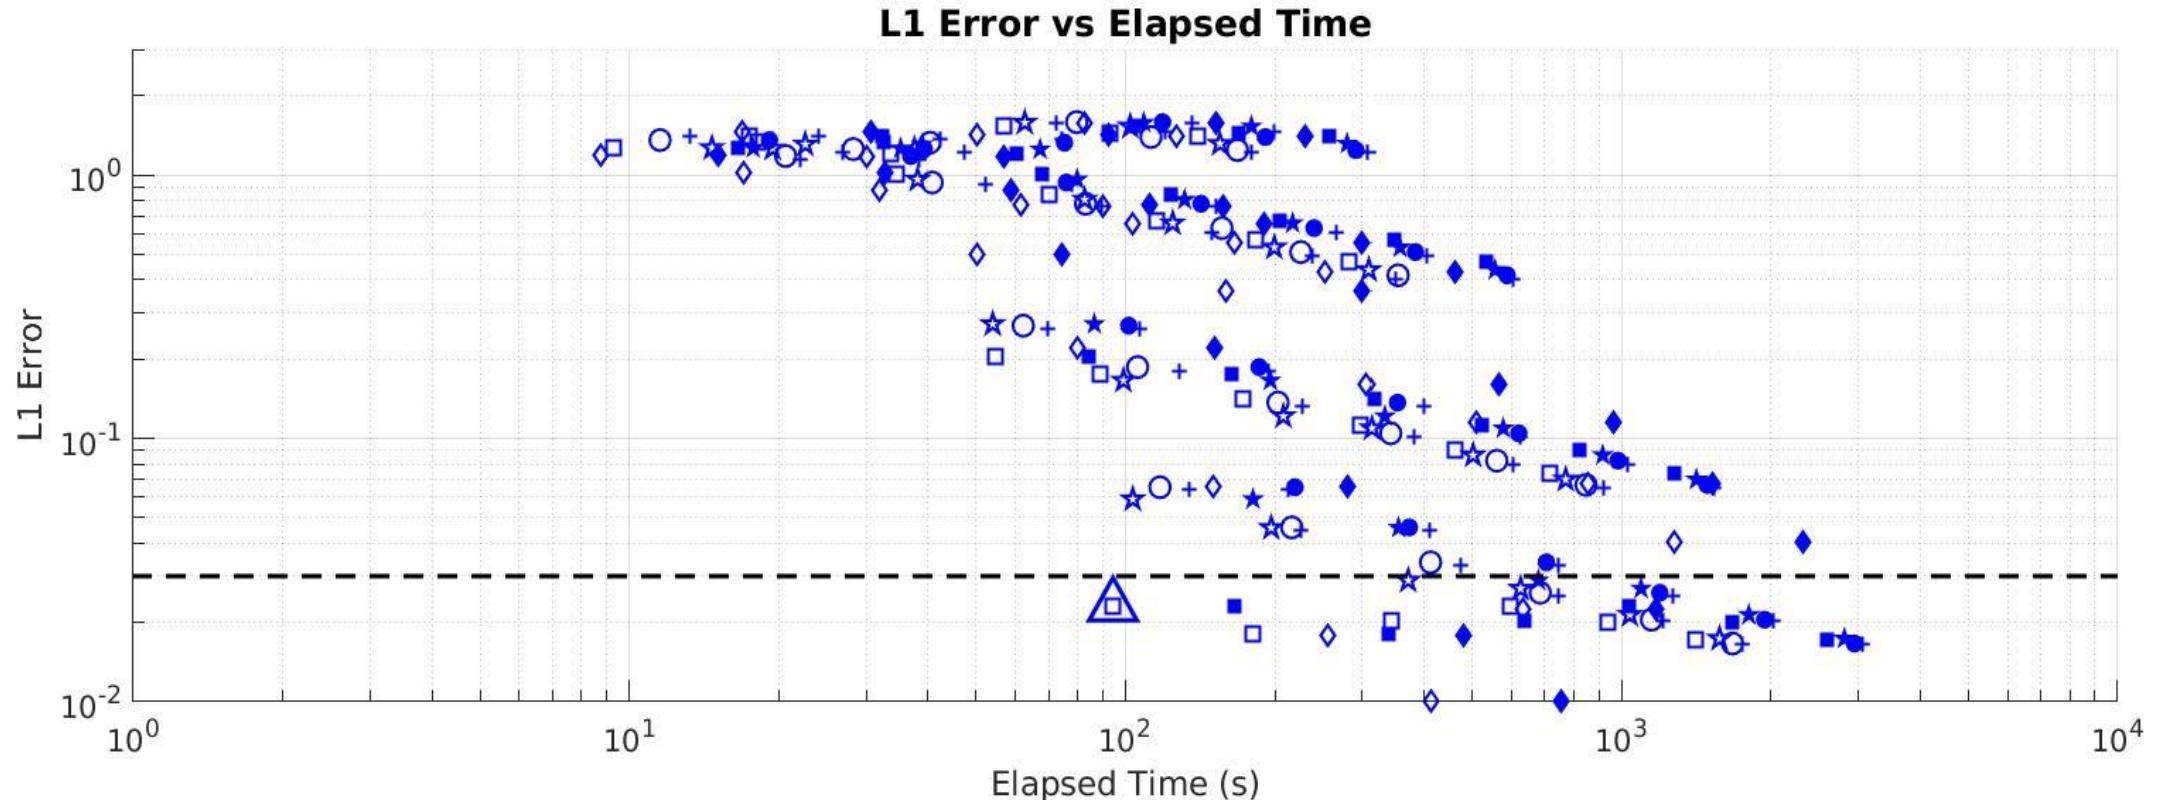
We scan range of parameters: stencil, space & time sampling

### **And we seek for the optimal configuration**

# Performance

Test case: Acoustic propagator  
Parametric analysis - Run time

**Max allowed error 3%**  
**Case RTM production mode**



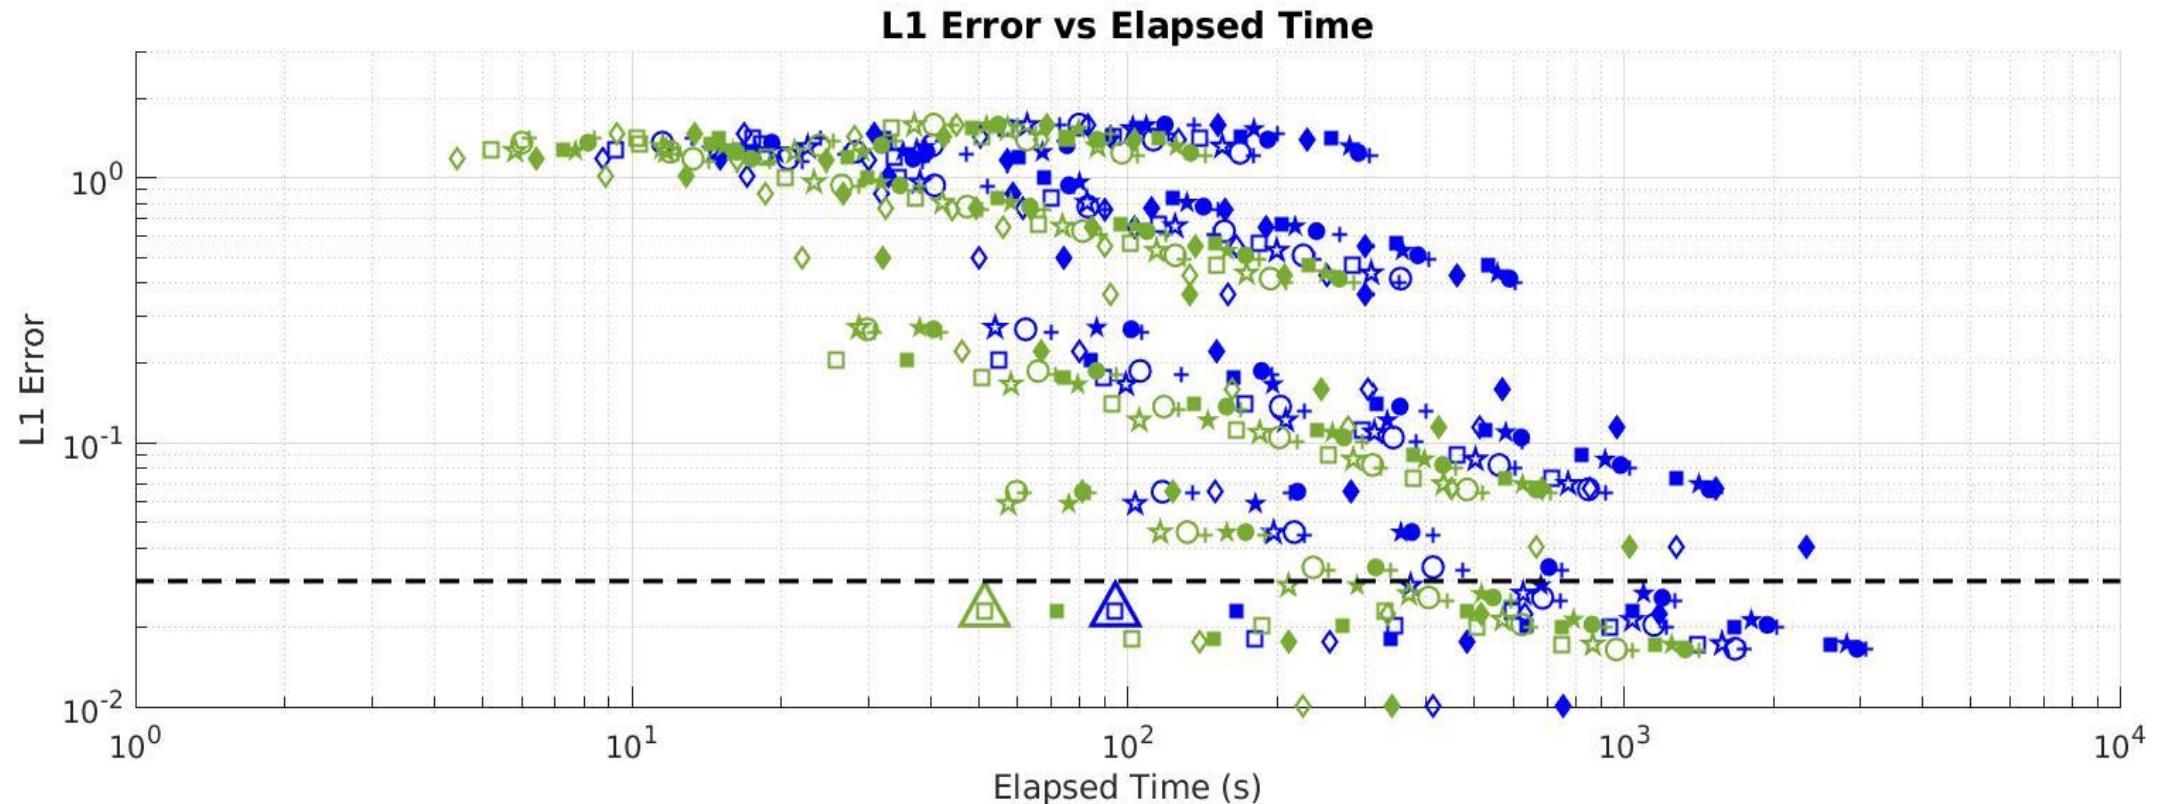
	CPU
Optimal algo.	Stand. O6
Optimal time (s)	94.4
Speed-up	1

Grid sizes: from 500x500x500 to 1000x1000x1000 (inc 100)  
FD order: from O4 to O12 (inc 2)  
CFL ratios: 1, 0.5, 0.2 & 0.1  
Propa. algo.: standard & split. computation  
**Total: 240 configurations**

# Performance

Test case: Acoustic propagator  
Parametric analysis - Run time

**Max allowed error 3%**  
**Case RTM production mode**

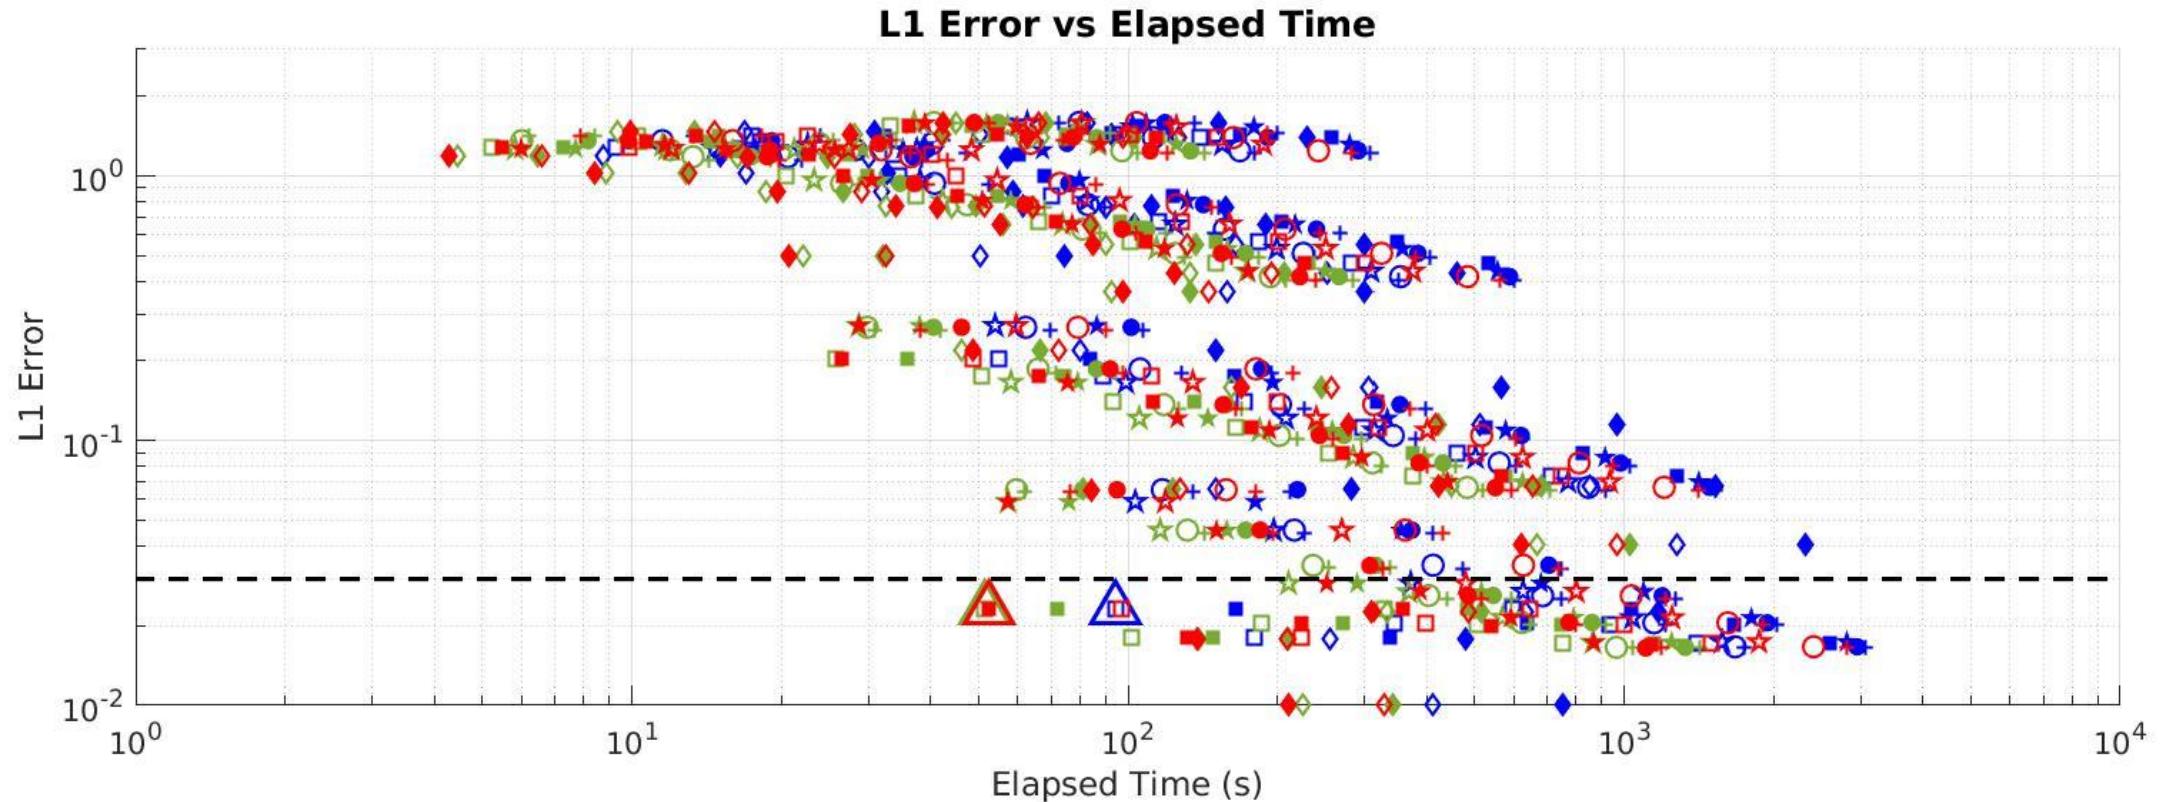


	CPU	GPU
Optimal algo.	Stand. O6	Stand. O6
Optimal time (s)	94.4	51.4
Speed-up	1	1.8

# Performance

Test case: Acoustic propagator  
Parametric analysis - Run time

**Max allowed error 3%**  
**Case RTM production mode**



	CPU	GPU	VE
Optimal algo.	Stand. O6	Stand. O6	Split O6
Optimal time (s)	94.4	51.4	52.5
Speed-up	1	1.8	1.8

**Optimal FD order 6**  
**GPU & VE: 2x speedup vs CPU**



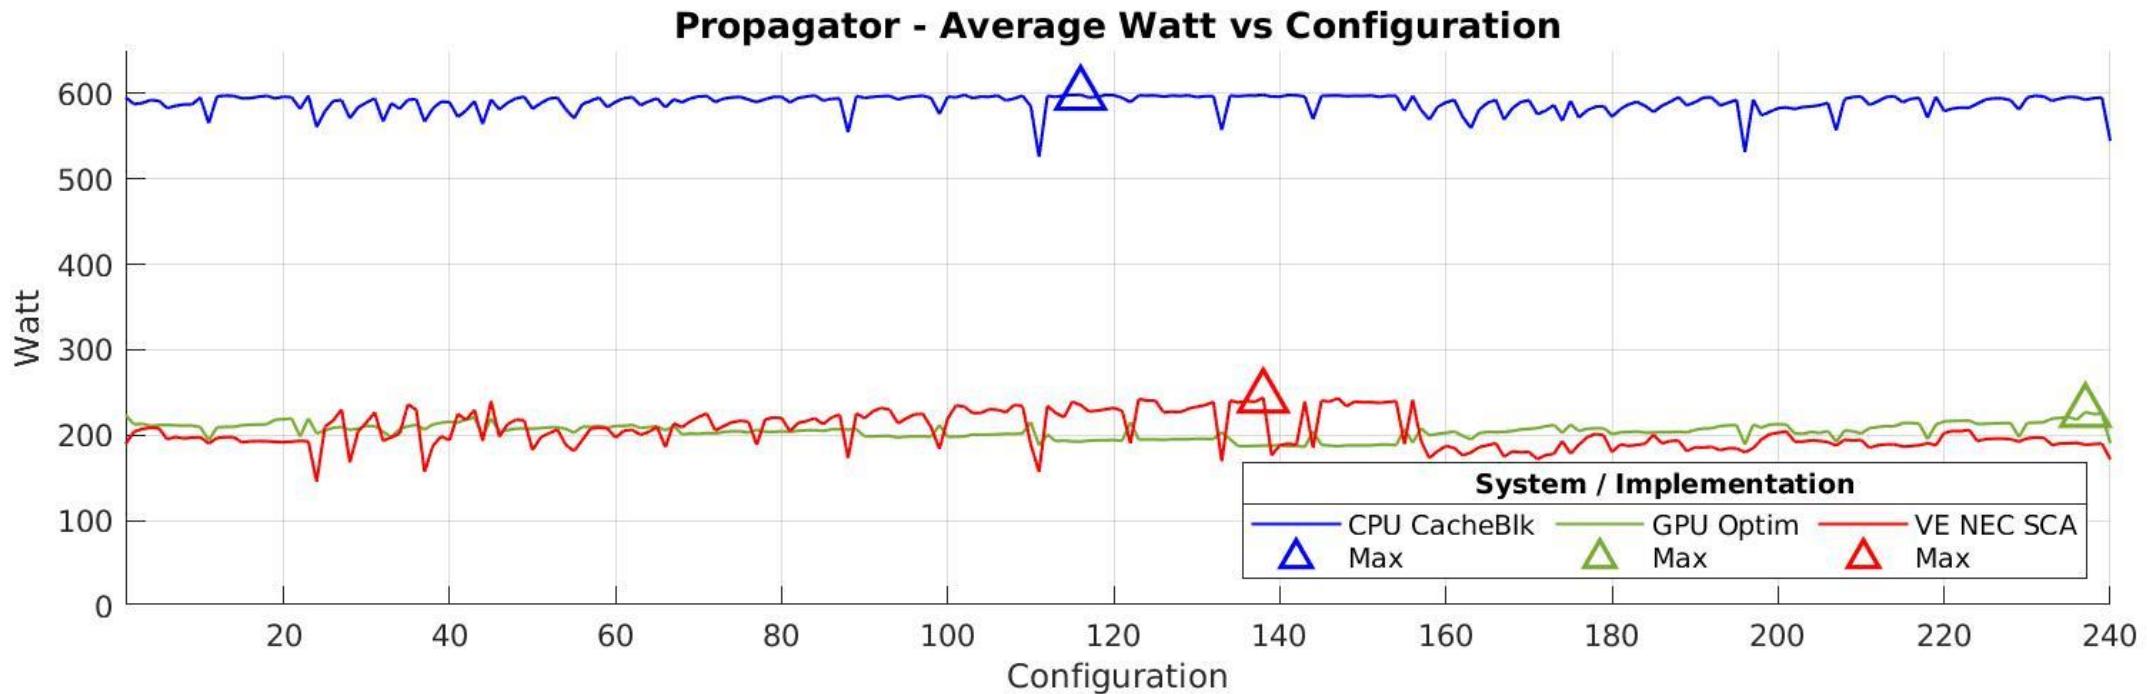
A characterization including  
an energy metrics is essential

$$\text{Watt} \times \text{runtime} = \text{W.h}$$

# Performance

## Test case: Acoustic propagator

### Parametric analysis - Energy



	CPU	GPU	VE
Peak Watt	600	300	300
Max measure	599	227	244
% peak	100	76	81

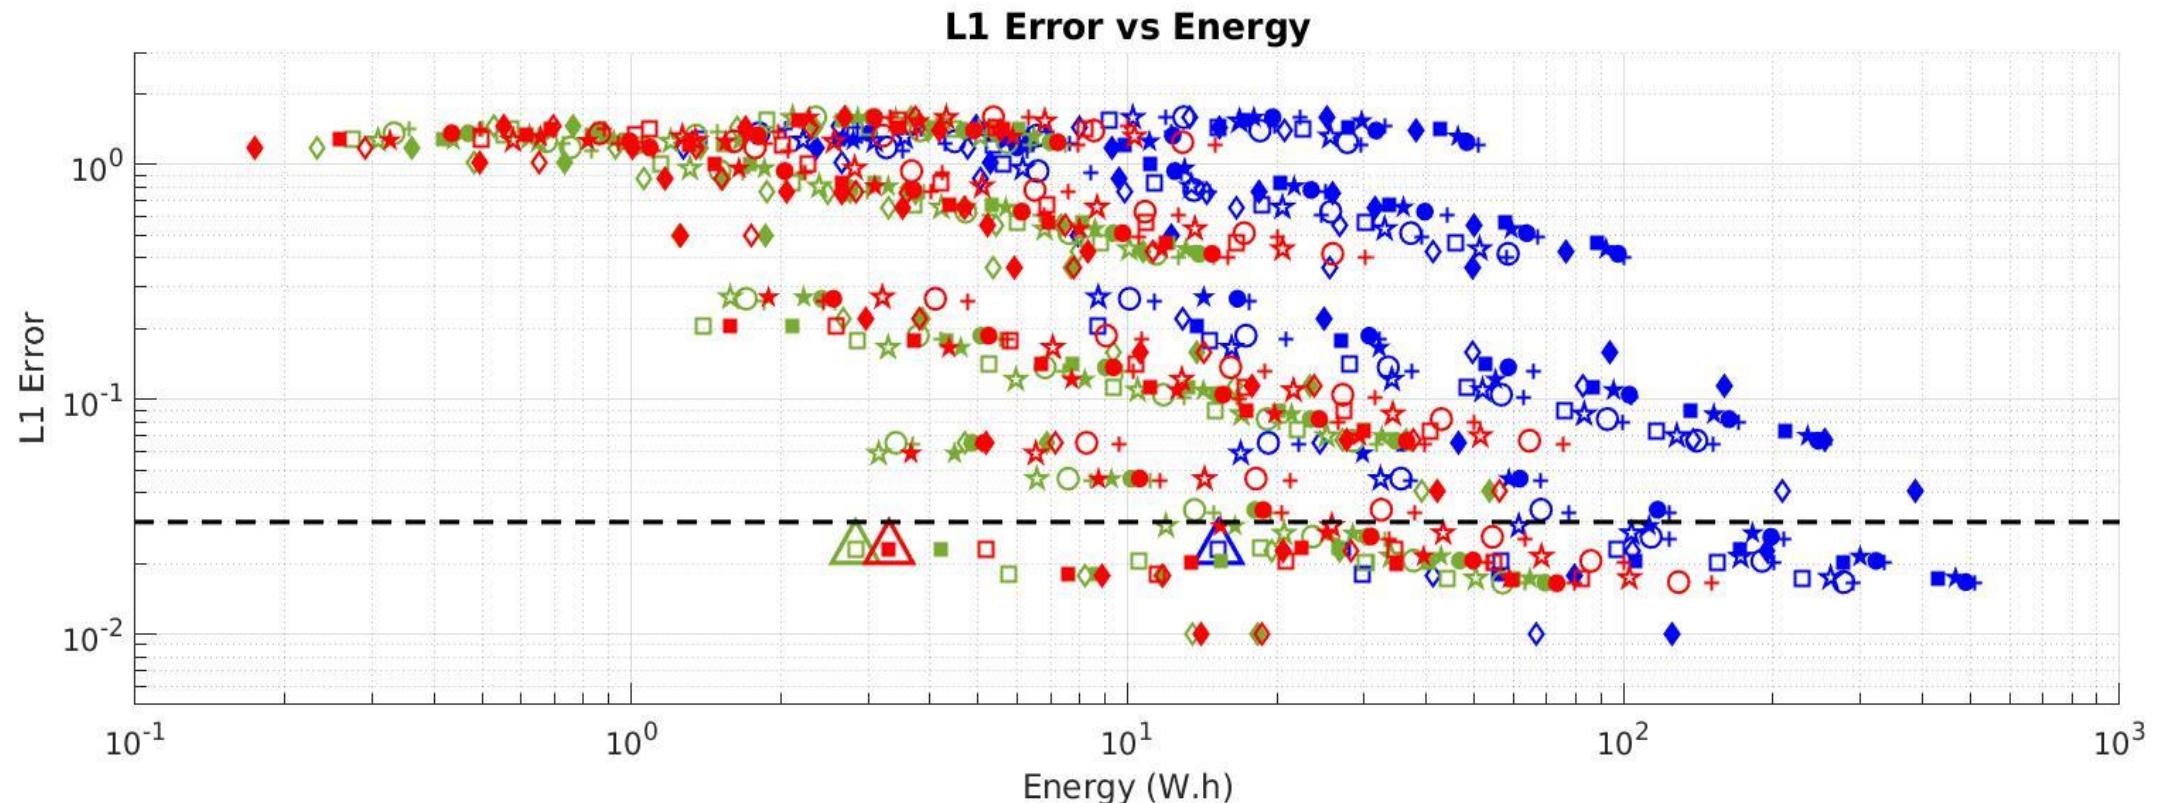
Algorithm	CPU	GPU	VE
Standard	585	208	190
Split Comp.	593	201	219

In average: GPU & VE: 2.5x less energy vs CPU

# Performance

## Test case: Acoustic propagator Parametric analysis - Energy

**Max allowed error 3%**  
**i.e. RTM production mode**



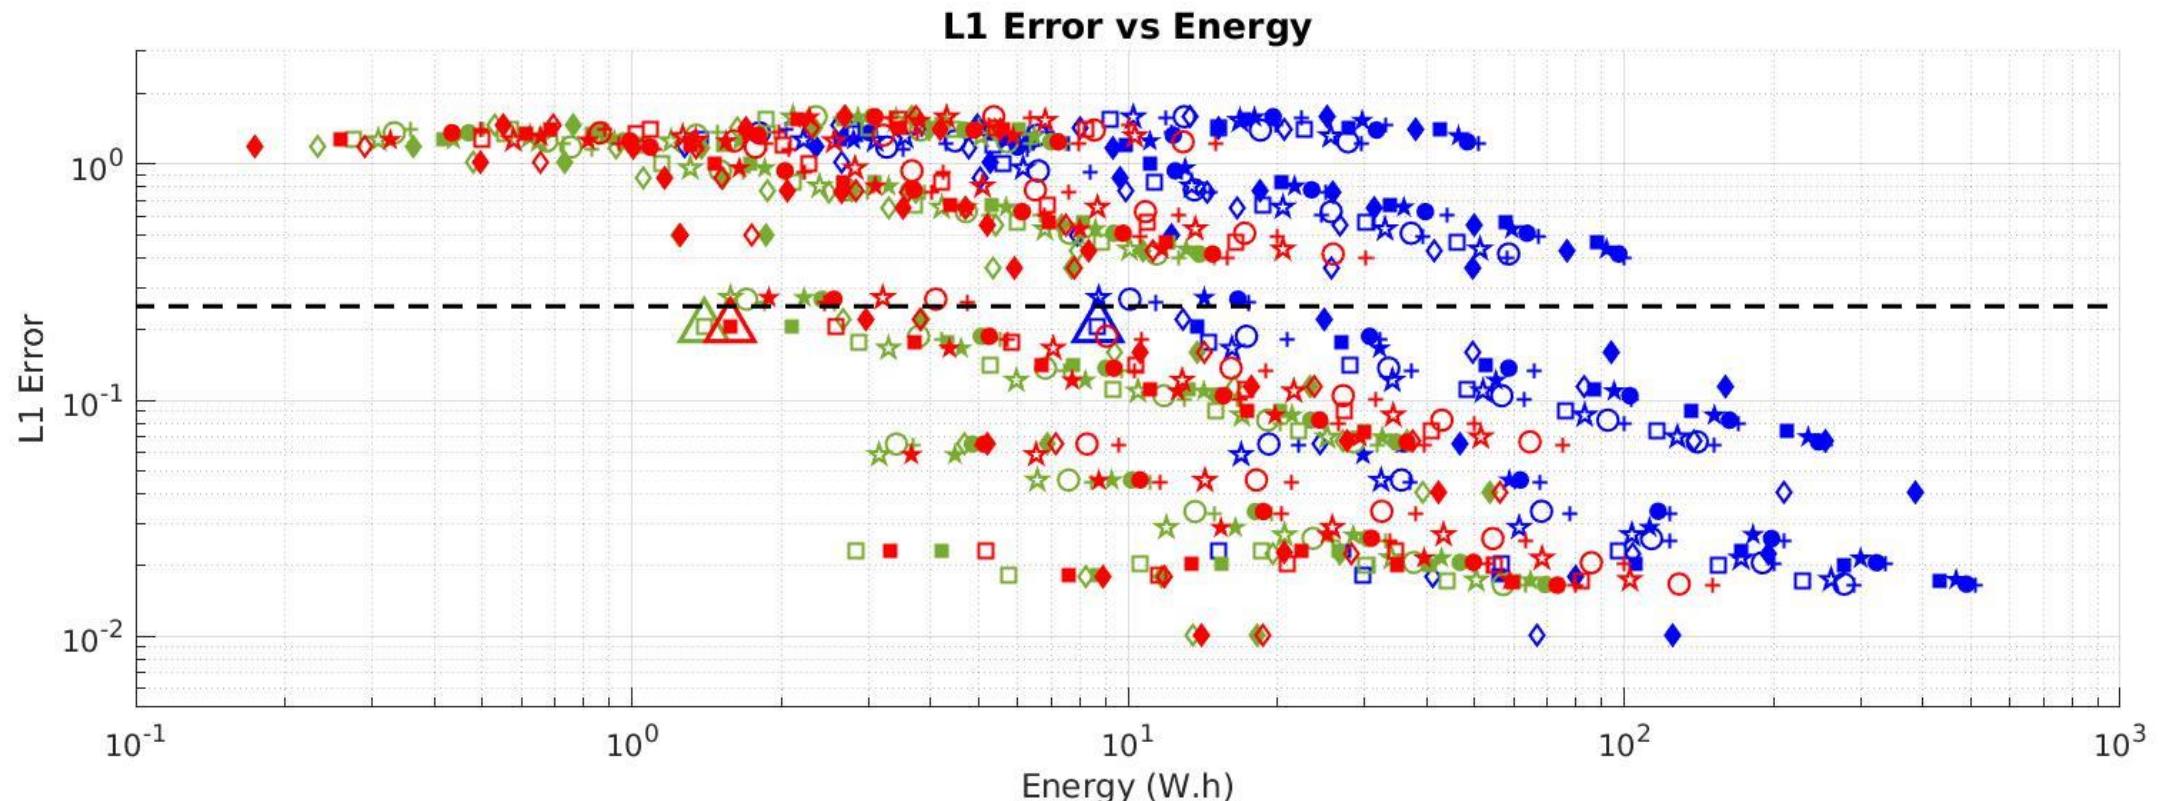
	CPU	GPU	VE
Optimal algo.	Stand. O6 N=500 0.1 CFL	Stand. O6 N=500 0.1 CFL	Split O6 N=500 0.1 CFL
Watt.Hour	15.24	2.83	3.31
Efficiency	1	5.4	4.6

**GPU & VE: Efficiency Combines Runtime + Energy**

# Performance

## Test case: Acoustic propagator Parametric analysis - Energy

**Max allowed error 25%** ↗  
i.e. RTM fast mode



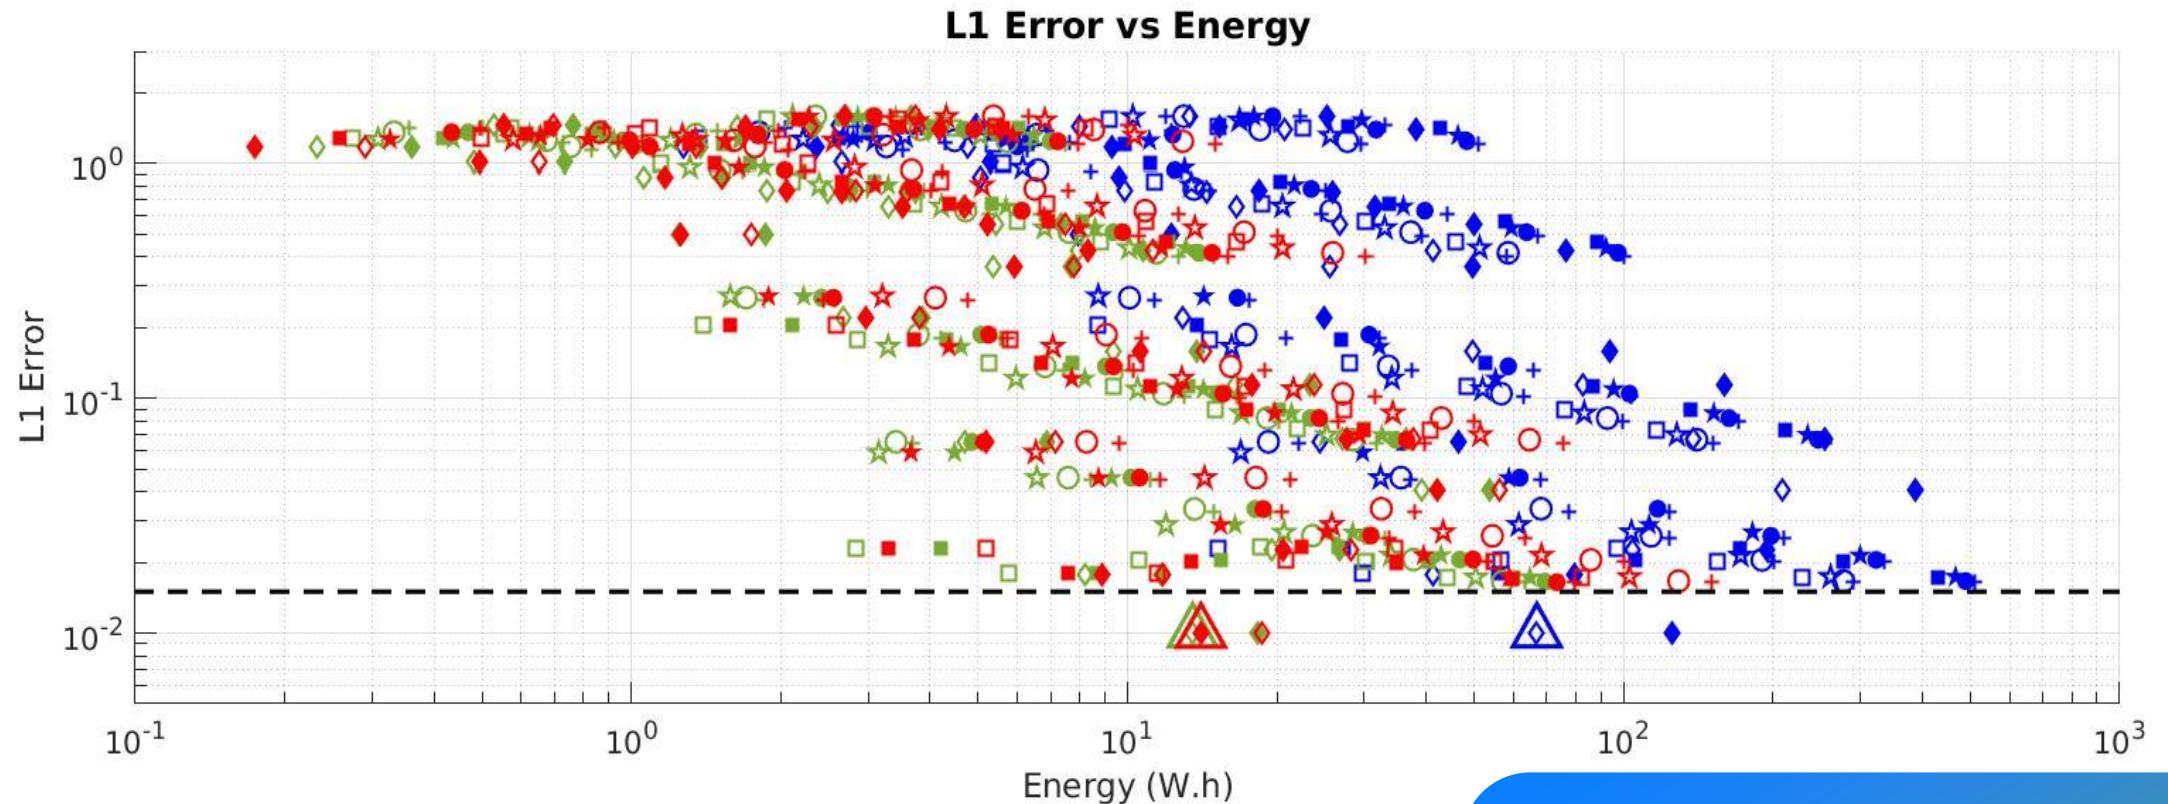
	CPU	GPU	VE
Optimal algo.	Stand. O6 N=500 0.2 CFL	Stand. O6 N=500 0.2 CFL	Split O6 N=500 0.2 CFL
Watt.Hour	8.70	1.40	1.58
Efficiency	1	6.2	5.5

To relax accuracy  
Changing  $\Delta t$  was most efficient

# Performance

## Test case: Acoustic propagator Parametric analysis - Energy

**Max allowed error 1.5%**  
i.e. modeling accurate mode



	CPU	GPU	VE
Optimal algo.	Stand. O4 N=900 0.2 CFL	Stand. O4 N=900 0.2 CFL	Split O4 N=900 0.2 CFL
Watt.Hour	66.8	13.5	14.1
Efficiency	1	5.0	4.7

To increase accuracy  
Smaller space sampling and  
lower order was most efficient  
However, ratio remains the  
same: GPU & VE 5x more  
efficient than CPU

# Conclusions

4

# Conclusions

We presented a **characterization of 3 HPC systems: CPU, GPU and VE** and showed

- the **importance of fine tuning for reliable comparison**
- the role of the **memory bandwidth and cache utilization** to achieve **high performance**

For an acoustic propagator (**compute intensive part of seismic imaging & inversion tools**)

**Performance of VE is similar to GPU**

- 2x runtime speed-up vs CPU
- **5x energy efficiency measured vs CPU**
- **Expected 4.4x energy efficiency at node level** (case of 8GPUs or 8VEs + 1 host (CPU) per node)

**High performance easily obtained with NEC SCA lib.** vs significant effort with GPU specific language

Future works

- Characterization of full application with all components (Interconnect + IO)
- More complex propagator (elastic vs acoustic)
- Higher order time integration scheme (high FD orders affected by 2<sup>nd</sup> order in time)
- Characterize more architectures (such as VE SX-Aurora 3, Q4 2022)

**You are welcome to join this initiative!**

# Acknowledgements



## Contributors to HPCscan

- Fabrice Dupros (Arm)
- Marcin Rogowski (KAUST)
- Philippe Thierry (Intel)
- Suha Kayum (Saudi Aramco)

## Discussions, Help and Support

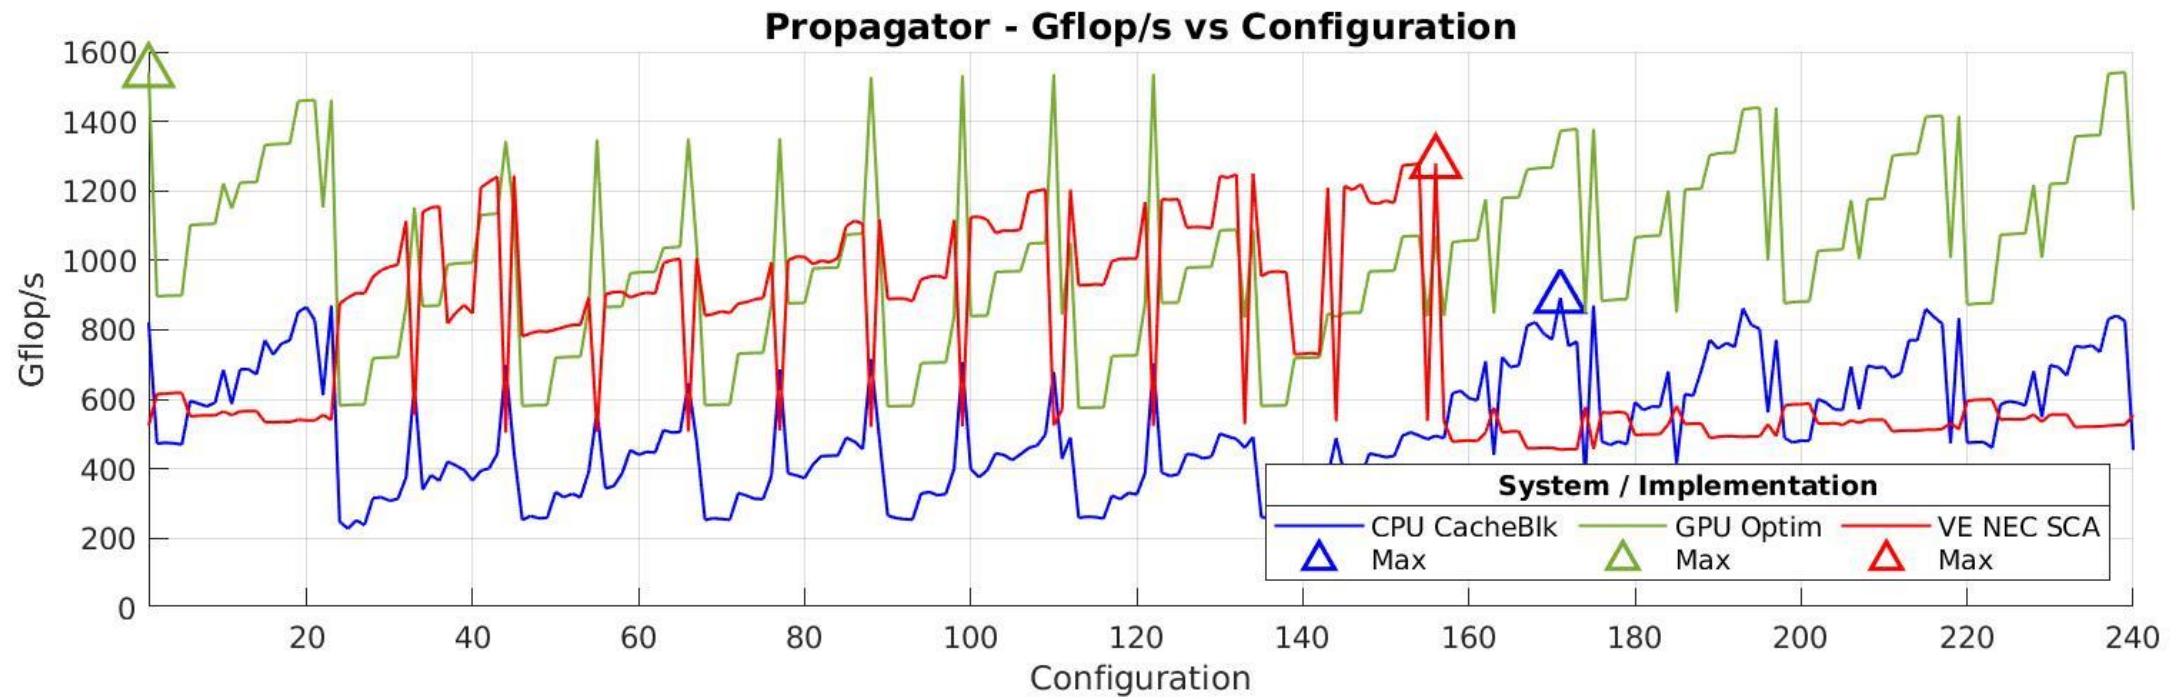
- Adil Lashab, Dominic Charrier, Mazda Sabony (AMD)
- Asma Farjallah, Rached Abdelkhalek (NVIDIA)
- Bruno Lecointe (NEC)
- Yoshio Nakao (Fujitsu)



# Backup Slides

# Performance

## Test case: Acoustic propagator Parametric analysis - Run time

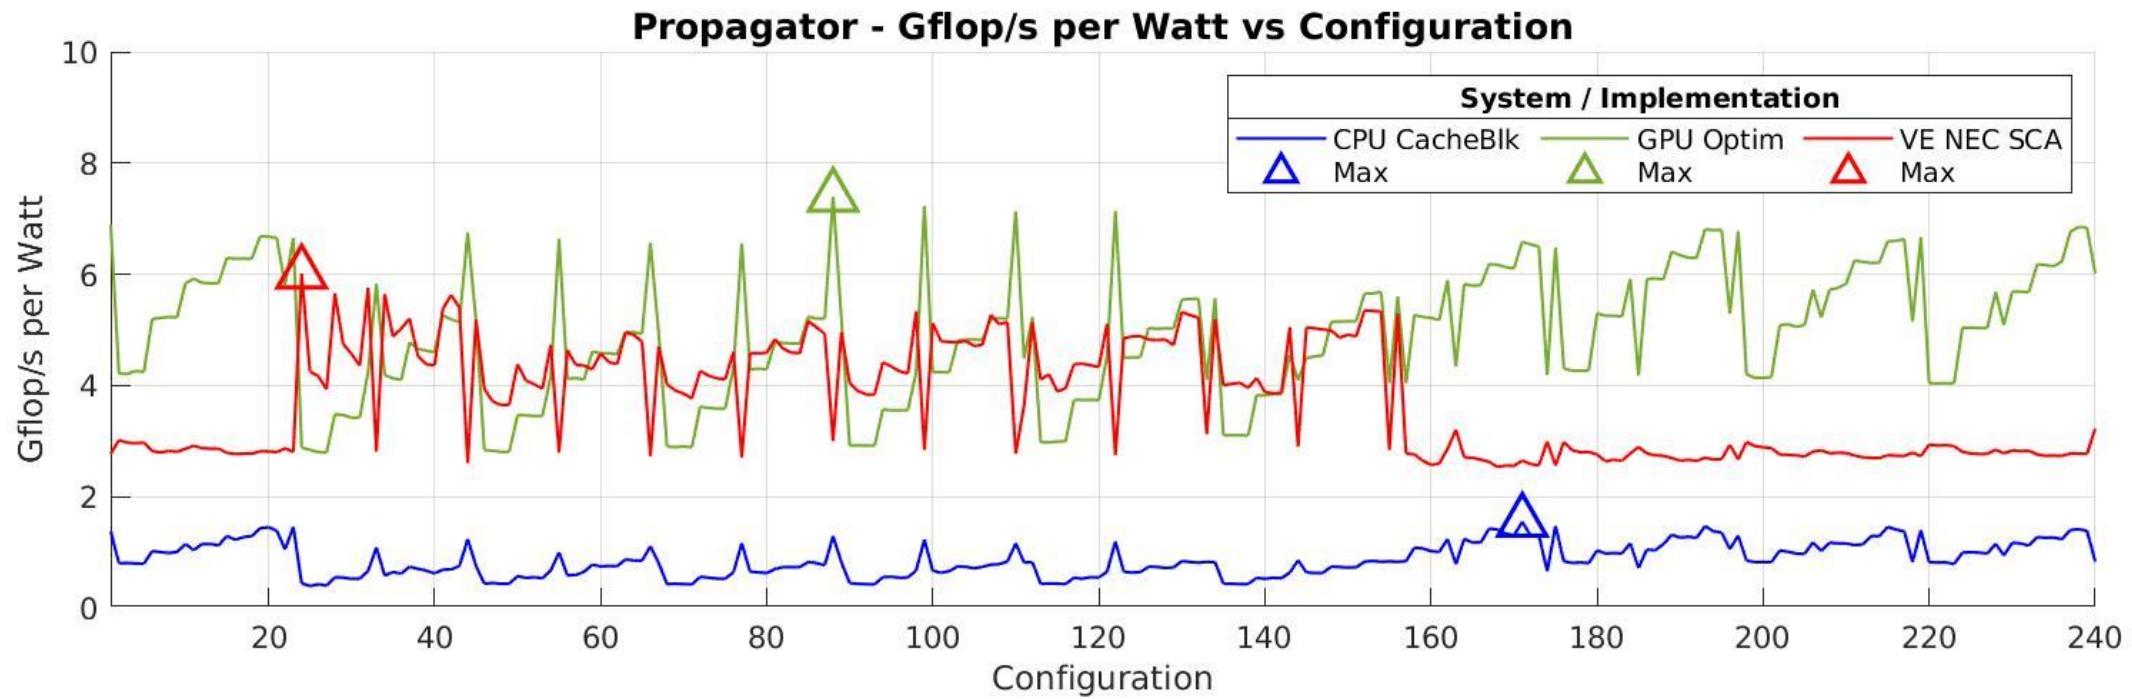


	CPU	GPU	VE
Peak Gflop/s	11 670	15 000	4 300
Max measure	892	1 542	1 280
% peak	7.6	10.3	29.8

Algorithm	CPU	GPU	VE
Standard	652	1179	531
Split Comp.	373	843	1 012

# Performance

## Test case: Acoustic propagator Parametric analysis - Energy



	CPU	GPU	VE
Peak Gflop/s / Watt	19.5	50	14.3
Max measure	1.5	7.4	6.0
% peak	7.7	14.8	42.0

Algorithm	CPU	GPU	VE
Standard	1.1	5.6	2.8
Split Comp.	0.6	4.2	4.6